

Aprendizado de Máquina e Reconhecimento de Padrões 2021.2



Evaluation Metrics for Classification

Prof. Samuel Martins (Samuka)
samuel.martins@ifsp.edu.br



Measuring Performance

When designing **a new model/solution** for a given problem, we desire to know about its **quality**: how **good** is your model? The common way to do that is to calculate some **evaluation metric** that **quantitatively** expresses the **model's quality**.

Measuring Performance

When designing **a new model/solution** for a given problem, we desire to know about its **quality**: how **good** is your model?

The common way to do that is to calculate some **evaluation metric** that **quantitatively** expresses the **model's quality**.

Performance

- How **good** is our model?
- How does it compare to others?

Measuring Performance

When designing **a new model/solution** for a given problem, we desire to know about its **quality**: how **good** is your model?

The common way to do that is to calculate some **evaluation metric** that **quantitatively** expresses the **model's quality**.

Performance

- How **good** is our model?
- How does it compare to others?

A good **evaluation metric** should:

- **Quantity** the **quality** of the model (output)
- Be directly comparable
- Do not be biased

Basic Definition

Positive

- The class/label of interest

Negative

- The other classes/labels.

Confusion Matrix

Confusion Matrix

		Predicted Label	
		Negative	Positive
Actual Label	Healthy	Healthy	Cancer
	Cancer		

The confusion matrix shows the performance of a diagnostic test. The columns represent the Predicted Label (Negative or Positive) and the rows represent the Actual Label (Healthy or Cancer). The matrix is color-coded: Yellow for Healthy-Predicted-Negative, Green for Healthy-Predicted-Positive, Red for Cancer-Predicted-Negative, and Teal for Cancer-Predicted-Positive.

		Predicted Label	
		Negative	Positive
Actual Label	Healthy	Healthy	Cancer
	Cancer		

Confusion Matrix

		Predicted Label	
		Negative	Positive
		Healthy	Cancer
Actual Label	Negative	Healthy	
	Positive	Cancer	

For simplification, we are considering 'healthy' as a person that **does not have cancer**. But note that this person may have other diseases, so that (s)he is not necessarily healthy.

Confusion Matrix

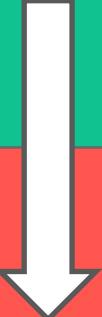
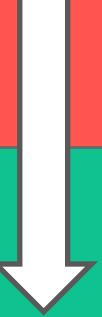
		Predicted Label	
		Negative	Positive
Actual Label	Healthy	Healthy	Cancer
	Cancer		

The confusion matrix illustrates the performance of a diagnostic test. The columns represent the Predicted Label (Negative or Positive) and the rows represent the Actual Label (Healthy or Cancer). The matrix shows the following results:

		Predicted Label	
		Negative	Positive
Actual Label	Healthy	True Negatives (Yellow)	False Positives (Red)
	Cancer	False Negatives (Red)	True Positives (Teal)

Confusion Matrix

the obtained prediction by
our model/solution

		Predicted Label	
		Negative	Positive
Actual Label	Healthy	Healthy	Cancer
	Cancer		

Confusion Matrix

known ground-truth

		Predicted Label	
		Negative	Positive
Actual Label	Healthy	Healthy	Cancer
	Healthy		
Positive	Cancer		

Confusion Matrix

		Predicted Label	
		Negative	Positive
Actual Label	Healthy	Healthy	Cancer
	Cancer		

		Predicted Label	
		Negative	Positive
Actual Label	Healthy	Healthy	Cancer
	Cancer		

Confusion Matrix

The **healthy patients** (those that **do not have cancer**) that **were correctly classified** by the model.

		Predicted Label	
		Negative	Positive
Actual Label	Negative	Healthy	Cancer
	Positive	True Negatives	
	Healthy		
	Cancer		

Confusion Matrix

		Predicted Label	
		Negative	Positive
Actual Label	Negative	Healthy	Cancer
	Positive	True Negatives	True Positives
Healthy			
Cancer			

The patients that **had cancer** that **were correctly classified** by the model.

Confusion Matrix

		Predicted Label	
		Negative	Positive
Actual Label	Negative	Healthy	Cancer
	Positive	True Negatives	False Positives
		Cancer	True Positives

Also called **Type I error** or **false alarm**

The **healthy patients** (those that **do not have cancer**) but the model **incorrectly** said **they have**.

Confusion Matrix

		Predicted Label	
		Negative	Positive
Actual Label	Negative	Healthy	Cancer
	Positive	True Negatives	False Positives
	Negative	False Negatives	True Positives
	Positive		

The cell containing "False Negatives" is highlighted with a black border.

A curved arrow points from the text below to this highlighted cell.

The patients that **have cancer** but the model **incorrectly** said they **did not**.

Also called **Type II error** or **miss**

Type I Error (False Positive)



Type II Error (False Negative)



Type I Error (False Positive)



Type II Error (False Negative)



Which is worse?

Type I Error (False Positive)



Type II Error (False Negative)



Which is worse?

It depends on the
problem/context/scenario.

You're not
pregnant

A scenario where...

A scenario where...

Type I Error is worse (False Positive)



An **innocent (negative)** on trial found
guilty (positive).

The judicial system seeks **to minimize Type I errors**
presumption of innocence

A scenario where...

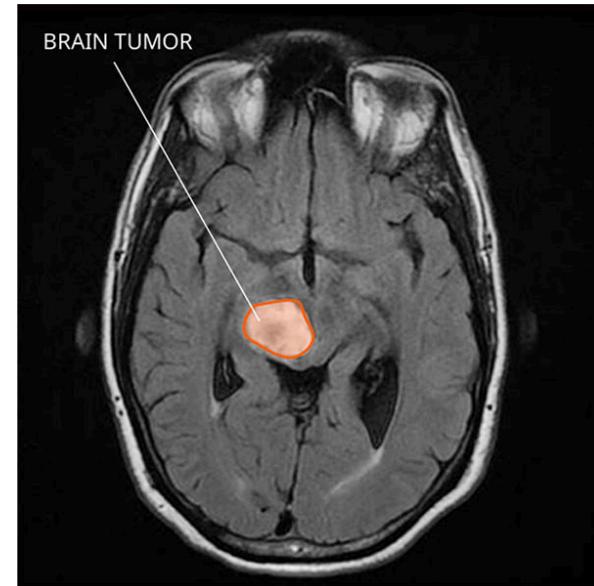
Type I Error is worse (False Positive)



An **innocent (negative)** on trial found guilty (**positive**).

The judicial system seeks **to minimize Type I errors**
presumption of innocence

Type II Error is worse (False Negative)



A medical screening incorrectly states that a **cancer patient (positive)** is **healthy (negative)**.

A scenario where...

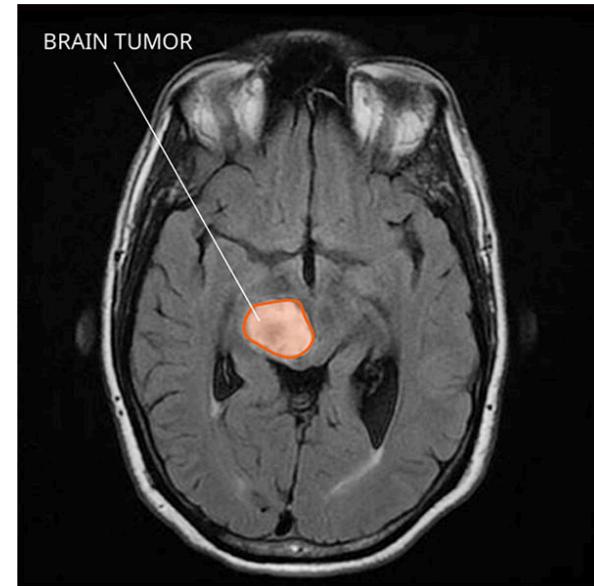
Type I Error is worse (False Positive)



An **innocent (negative)** on trial found guilty (**positive**).

The judicial system seeks **to minimize Type I errors**
presumption of innocence

Type II Error is worse (False Negative)



A medical screening incorrectly states that a **cancer patient (positive)** is **healthy (negative)**.



Type II error seems to be **worse** for most problems.

Confusion Matrix

		Predicted Label	
		Negative	Positive
Actual Label	Negative	Healthy	Cancer
	Positive	True Negatives	False Positives
	Negative	False Negatives	True Positives
	Positive	Cancer	

Confusion Matrix

		Predicted Label	
		Negative	Positive
Actual Label	Healthy	Healthy	Cancer
	Healthy	TN	FP
Positive	Cancer	FN	TP

Confusion Matrix

		Predicted Label	
		Negative	Positive
Actual Label	Healthy	Healthy	Cancer
	Healthy	TN	FP
Positive	Cancer	FN	TP
	Cancer		



You may see this matrix with columns and rows in a different order.

Example

A given **model's predictions** for a Validation set with:

- 100 people that **have cancer**
- 200 people that **do not have cancer (healthy)**.

		Predicted Label	
		Negative	Positive
		Healthy	Cancer
Actual Label	Negative	Healthy	TN
	Positive	Cancer	FP FN TP

Example

A given **model's predictions** for a Validation set with:

- 100 people that **have cancer**
- 200 people that **do not have cancer (healthy)**.

From **200 people that do not have cancer**:

		Predicted Label	
		Negative	Positive
Actual Label	Negative	Healthy	Cancer
	Healthy	TN	FP
Positive	Cancer	FN	TP

Example

A given **model's predictions** for a Validation set with:

- 100 people that **have cancer**
- 200 people that **do not have cancer (healthy)**.

From **200 people that do not have cancer**:

- **70** were **correctly classified**

		Predicted Label	
		Negative	Positive
Actual Label	Negative	Healthy	Cancer
	Positive	70	FP
Healthy	Cancer	FN	TP

Example

A given **model's predictions** for a Validation set with:

- 100 people that **have cancer**
- 200 people that **do not have cancer (healthy)**.

		Predicted Label	
		Negative	Positive
Actual Label	Negative	Healthy	Cancer
	Positive	Healthy	70
		Cancer	FN
			TP

From **200 people that do not have cancer**:

- **70** were **correctly classified**
- **130** were **incorrectly classified** were as **having cancer**.

Example

A given **model's predictions** for a Validation set with:

- 100 people that **have cancer**
- 200 people that **do not have cancer (healthy)**.

		Predicted Label	
		Negative	Positive
Actual Label	Negative	Healthy	Cancer
	Healthy	70	130
Positive	Cancer	FN	TP

From **100 people that have cancer**:

Example

A given **model's predictions** for a Validation set with:

- 100 people that **have cancer**
- 200 people that **do not have cancer (healthy)**.

		Predicted Label	
		Negative	Positive
Actual Label	Negative	Healthy	Cancer
	Positive	Healthy	70
		Cancer	FN
			80

From 100 people that have cancer:
- 80 were **correctly classified**

Example

A given **model's predictions** for a Validation set with:

- 100 people that **have cancer**
- 200 people that **do not have cancer (healthy)**.

		Predicted Label	
		Negative	Positive
Actual Label	Negative	Healthy	Cancer
	Positive	70	130
		20	80

From **100 people that have cancer**:

- 80 were **correctly classified**
- 20 were **incorrectly classified** as **healthy**.

Confusion Matrix

```
sklearn.metrics.confusion_matrix
```

```
from sklearn.metrics import confusion_matrix  
  
tn, fp, fn, tp = confusion_matrix(y_true, y_pred).ravel()
```

		Predicted Label	
		0	1
Actual Label	0	TN	FP
	1	FN	TP

Confusion Matrix

```
sklearn.metrics.confusion_matrix
```

```
from sklearn.metrics import confusion_matrix  
  
tn, fp, fn, tp = confusion_matrix(y_true, y_pred).ravel()
```

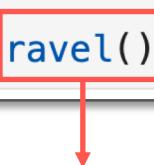
For **binary** problems, it returns a **2x2 matrix** (numpy array)

		Predicted Label	
		0	1
Actual Label	0	TN	FP
	1	FN	TP

Confusion Matrix

```
sklearn.metrics.confusion_matrix
```

```
from sklearn.metrics import confusion_matrix  
  
tn, fp, fn, tp = confusion_matrix(y_true, y_pred).ravel()
```



It transforms the **2x2 matrix** into a **1D numpy** array by flattening it.

		Predicted Label	
		0	1
Actual Label	0	TN	FP
	1	FN	TP

Confusion Matrix

```
sklearn.metrics.confusion_matrix
```

```
from sklearn.metrics import confusion_matrix  
  
tn, fp, fn, tp = confusion_matrix(y_true, y_pred).ravel()
```

⚠ The classes can be represented with any value (string or int), not necessarily 0 and 1.

⚠ If the problem is **binary**, the **negative class** will be that **with the lowest value**.

		Predicted Label	
		0	1
Actual Label	0	TN	FP
	1	FN	TP

Confusion Matrix

`sklearn.metrics.confusion_matrix`

```
from sklearn.metrics import confusion_matrix  
  
tn, fp, fn, tp = confusion_matrix(y_true, y_pred).ravel()
```

In case of **strings**, the function will consider as **negative** the one that comes **alphabetically first**.
So, in our problem, **Cancer** would be considered the **negative** class.

		Predicted Label	
		0	1
Actual Label	0	TN	FP
	1	FN	TP

Confusion Matrix

sklearn.metrics.confusion_matrix

```
from sklearn.metrics import confusion_matrix  
  
tn, fp, fn, tp = confusion_matrix(y_true, y_pred,  
                                    labels=['Healthy', 'Cancer']).ravel()
```

Negative Positive

! In case of **strings**, the function will consider as **negative** the one that comes **alphabetically first**. So, in our problem, **Cancer** would be considered the **negative** class.



We can **define** which class will be the **negative** and **positive** classes by passing the order in **'labels'**.

		Predicted Label	
		Healthy	Cancer
Actual Label	Healthy	TN	FP
	Cancer	FN	TP

Comparing Confusion Matrices

Confusion matrix for classifier A (e.g., KNN)

		Predicted Label	
		Negative	Positive
		Healthy	Cancer
Actual Label Positive	Healthy	70	130
	Cancer	20	80

Confusion matrix for classifier A
(e.g., KNN)

		Predicted Label	
		Negative	Positive
		Healthy	Cancer
Actual Label Positive	Healthy	70	130
	Cancer	20	80

Confusion matrix for classifier B
(e.g., Logistic Regression)

		Predicted Label	
		Negative	Positive
		Healthy	Cancer
Actual Label Positive	Healthy	80	120
	Cancer	50	50

Confusion matrix for classifier A
(e.g., KNN)

		Predicted Label	
		Negative	Positive
		Healthy	Cancer
Actual Label	Positive	70	130
Positive	Negative	20	80

Confusion matrix for classifier B
(e.g., Logistic Regression)

		Predicted Label	
		Negative	Positive
		Healthy	Cancer
Actual Label	Positive	80	120
Positive	Negative	50	50

KNN was better than Logistic Regression at classifying patients with cancer (80 vs 50)...

Confusion matrix for classifier A
(e.g., KNN)

		Predicted Label	
		Negative	Positive
Actual Label	Healthy	70	130
	Cancer	20	80

Confusion matrix for classifier B
(e.g., Logistic Regression)

		Predicted Label	
		Negative	Positive
Actual Label	Healthy	80	120
	Cancer	50	50

... but it is worse at classifying **healthy patients (70 vs 80)**.

Confusion matrix for classifier A
(e.g., KNN)

		Predicted Label	
		Negative	Positive
		Healthy	Cancer
Actual Label Positive Negative	Healthy	70	130
	Cancer	20	80

Confusion matrix for classifier B
(e.g., Logistic Regression)

		Predicted Label	
		Negative	Positive
		Healthy	Cancer
Actual Label Positive Negative	Healthy	80	120
	Cancer	50	50

So, if it is **more important** to **correctly classify positives**, we should choose **KNN**

Confusion matrix for classifier A
(e.g., KNN)

		Predicted Label	
		Negative	Positive
		Healthy	Cancer
Actual Label Positive	Healthy	70	130
	Cancer	20	80

Confusion matrix for classifier B
(e.g., Logistic Regression)

		Predicted Label	
		Negative	Positive
		Healthy	Cancer
Actual Label Positive	Healthy	80	120
	Cancer	50	50

Otherwise, if it is **more important to correctly classify negatives**, we should choose **Logistic Regression**

Confusion matrix for classifier A
(e.g., KNN)

		Predicted Label	
		Negative	Positive
		Healthy	Cancer
Actual Label Positive	Healthy	70	130
	Cancer	20	80

Confusion matrix for classifier B
(e.g., Logistic Regression)

		Predicted Label	
		Negative	Positive
		Healthy	Cancer
Actual Label Positive	Healthy	80	120
	Cancer	50	50

Comparing **confusion matrices** can be **hard** especially when they are **very similar**, so making it difficult to know which model is better,

Confusion matrix for classifier A
(e.g., KNN)

		Predicted Label	
		Negative	Positive
		Healthy	Cancer
Actual Label Positive	Healthy	70	130
	Cancer	20	80

Comparing **confusion matrices** can be **hard** especially when they are **very similar**, so making it difficult to know which model is better,



Confusion matrix for classifier B
(e.g., Logistic Regression)

		Predicted Label	
		Negative	Positive
		Healthy	Cancer
Actual Label Positive	Healthy	80	120
	Cancer	50	50

There are more sophisticated **evaluation metrics**, such as **Sensitivity**, **Specificity**, **Precision**, **Recall**, ..., that can help us make a better decision by **summarizing** some points of the **confusion matrix**.



Confusion Matrix for Multiple Classes

		Predicted Label		
		Cat	Dog	Wolf
Actual Label	Cat	TP _{Cat}		
	Dog		TP _{Dog}	
	Wolf			TP _{Wolf}

Confusion Matrix for Multiple Classes

The **diagonal** keeps showing where the model/solution **hit**, i.e., **correctly classified the classes (true positives)**.

		Predicted Label		
		Cat	Dog	Wolf
Actual Label	Cat	TP_{Cat}		
	Dog		TP_{Dog}	
	Wolf			TP_{Wolf}

Confusion Matrix for Multiple Classes

The **diagonal** keeps showing where the model/solution **hit**, i.e., **correctly classified the classes (true positives)**.

```
from sklearn.metrics import confusion_matrix  
confusion_matrix(y_true, y_pred, labels=['Cat', 'Dog', 'Wolf'])
```

		Predicted Label		
		Cat	Dog	Wolf
Actual Label	Cat	TP _{Cat}		
	Dog		TP _{Dog}	
	Wolf			TP _{Wolf}

Confusion Matrix for Multiple Classes

The **diagonal** keeps showing where the model/solution **hit**, i.e., **correctly classified the classes (true positives)**.

```
from sklearn.metrics import confusion_matrix  
confusion_matrix(y_true, y_pred, labels=['Cat', 'Dog', 'Wolf'])
```

Just to guarantee the **expected order** for the classes (rows and columns).



Otherwise, the **order** will be the **alphabetical one**, which is the same in this case.

		Predicted Label		
		Cat	Dog	Wolf
Actual Label	Cat	TP _{Cat}		
	Dog		TP _{Dog}	
	Wolf			TP _{Wolf}

Confusion Matrix for Multiple Classes

		Predicted Label		
		Cat	Dog	Wolf
Actual Label	Cat	10	8	2
	Dog	5	30	15
	Wolf	1	4	15

Confusion Matrix for Multiple Classes

The **diagonal** keeps showing where the model/solution **hit**, i.e., **correctly classified the classes (true positives)**.

		Predicted Label		
		Cat	Dog	Wolf
Actual Label	Cat	10	8	2
	Dog	5	30	15
	Wolf	1	4	15

Confusion Matrix for Multiple Classes

The remaining cells shows where the model/solution **missed**, i.e., **incorrectly classified the classes**.

		Predicted Label		
		Cat	Dog	Wolf
Actual Label	Cat	10	8	2
	Dog	5	30	15
	Wolf	1	4	15

Confusion Matrix for Multiple Classes

TP_{Cat}

		Predicted Label		
		Cat	Dog	Wolf
Actual Label	Cat	10	8	2
	Dog	5	30	15
	Wolf	1	4	15

Confusion Matrix for Multiple Classes

TP_{Cat} FP_{Cat}

		Predicted Label		
		Cat	Dog	Wolf
Actual Label	Cat	10	8	2
	Dog	5	30	15
	Wolf	1	4	15

Confusion Matrix for Multiple Classes

TP_{Cat}

FP_{Cat}

FN_{Cat}

		Predicted Label		
		Cat	Dog	Wolf
Actual Label	Cat	10	8	2
	Dog	5	30	15
	Wolf	1	4	15

Confusion Matrix for Multiple Classes

$$TP_{Cat} \quad FP_{Cat}$$
$$FN_{Cat} \quad TN_{Cat}$$

		Predicted Label		
		Cat	Dog	Wolf
Actual Label	Cat	10	8	2
	Dog	5	30	15
	Wolf	1	4	15

Confusion Matrix for Multiple Classes

TP_{Dog}

		Predicted Label		
		Cat	Dog	Wolf
Actual Label	Cat	10	8	2
	Dog	5	30	15
	Wolf	1	4	15

Confusion Matrix for Multiple Classes

TP_{Dog} FP_{Dog}

		Predicted Label		
		Cat	Dog	Wolf
Actual Label	Cat	10	8	2
	Dog	5	30	15
	Wolf	1	4	15

Confusion Matrix for Multiple Classes

TP_{Dog} FP_{Dog}
 FN_{Dog}

		Predicted Label		
		Cat	Dog	Wolf
Actual Label	Cat	10	8	2
	Dog	5	30	15
	Wolf	1	4	15

Confusion Matrix for Multiple Classes

$$\begin{matrix} \textcolor{blue}{TP}_{Dog} & \textcolor{purple}{FP}_{Dog} \\ \textcolor{black}{FN}_{Dog} & \textcolor{orange}{TN}_{Dog} \end{matrix}$$

		Predicted Label		
		Cat	Dog	Wolf
Actual Label	Cat	10	8	2
	Dog	5	30	15
	Wolf	1	4	15

Confusion Matrix for Multiple Classes

Exercise: Compute...

$$\begin{array}{ll} \textcolor{blue}{TP}_{\text{Wolf}} & \textcolor{purple}{FP}_{\text{Wolf}} \\ \textcolor{black}{FN}_{\text{Wolf}} & \textcolor{orange}{TN}_{\text{Wolf}} \end{array}$$

		Predicted Label		
		Cat	Dog	Wolf
Actual Label	Cat	10	8	2
	Dog	5	30	15
	Wolf	1	4	15

Confusion Matrix for Multiple Classes

If we had **c classes**, we would have a **confusion matrix** with **c rows and c columns**.

		Predicted Label					
		1	2	3	4	5	6
Actual Label	1	1	0	0	0	0	0
	2	0	1	0	0	0	0
3	1	0	0	1	0	0	0
	2	0	0	0	1	0	0
4	1	0	0	0	0	1	0
	2	0	0	0	0	0	1
5	1	0	0	0	0	0	1
	2	0	0	0	0	1	0
6	1	0	0	0	0	0	1
	2	0	0	0	0	1	0

Confusion Matrix for Multiple Classes

If we had **c classes**, we would have a **confusion matrix** with **c rows and c columns..**

		Predicted Label					
		TN		FP		TN	
Actual Label	TN	TP	FN	FP	TP	FN	
	TN	FP	TN	FP	TN		

For a given **class**

Confusion Matrix for Multiple Classes

```
from sklearn.metrics import confusion_matrix
import numpy as np

cm = confusion_matrix(y_true, y_pred,
                      labels=['Cat', 'Dog', 'Wolf'])

TP = np.diag(cm)
FP = cm.sum(axis=0) - TP
FN = cm.sum(axis=1) - TP
TN = cm.sum() - (TP + FP + FN)
```

Confusion Matrix for Multiple Classes

```
from sklearn.metrics import confusion_matrix
import numpy as np

cm = confusion_matrix(y_true, y_pred,
                      labels=['Cat', 'Dog', 'Wolf'])

TP = np.diag(cm)
FP = cm.sum(axis=0) - TP
FN = cm.sum(axis=1) - TP
TN = cm.sum() - (TP + FP + FN)
```

		Predicted Label		
		0	1	2
Actual Label	0	10	8	2
	1	5	30	15
	2	1	4	15

Confusion Matrix for Multiple Classes

```
from sklearn.metrics import confusion_matrix
import numpy as np

cm = confusion_matrix(y_true, y_pred,
                      labels=['Cat', 'Dog', 'Wolf'])

TP = np.diag(cm)
FP = cm.sum(axis=0) - TP
FN = cm.sum(axis=1) - TP
TN = cm.sum() - (TP + FP + FN)
```

0: Cat
1: Dog
2: Wolf

		Predicted Label		
		0	1	2
Actual Label	0	10	8	2
	1	5	30	15
	2	1	4	15

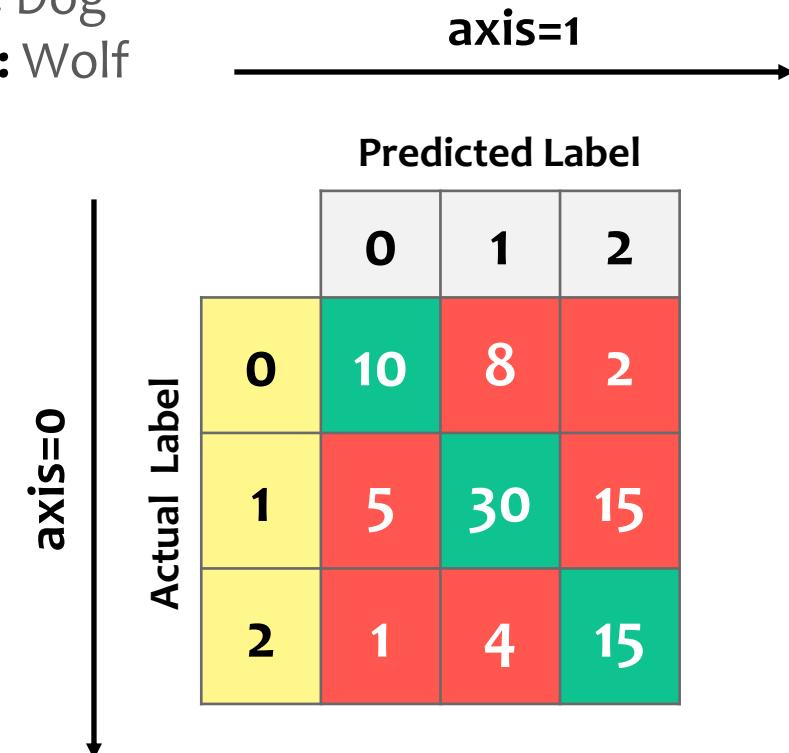
Confusion Matrix for Multiple Classes

```
from sklearn.metrics import confusion_matrix
import numpy as np

cm = confusion_matrix(y_true, y_pred,
                      labels=['Cat', 'Dog', 'Wolf'])

TP = np.diag(cm)
FP = cm.sum(axis=0) - TP
FN = cm.sum(axis=1) - TP
TN = cm.sum() - (TP + FP + FN)
```

0: Cat
1: Dog
2: Wolf



Confusion Matrix for Multiple Classes

```
from sklearn.metrics import confusion_matrix
import numpy as np

cm = confusion_matrix(y_true, y_pred,
                      labels=['Cat', 'Dog', 'Wolf'])

TP = np.diag(cm)
FP = cm.sum(axis=0) - TP
FN = cm.sum(axis=1) - TP
TN = cm.sum() - (TP + FP + FN)
```

$$TP = \begin{array}{|c|c|c|} \hline 0 & 1 & 2 \\ \hline \end{array}$$

0: Cat
1: Dog
2: Wolf

axis=1

		Predicted Label		
		0	1	2
Actual Label	0	10	8	2
	1	5	30	15
2	1	4	15	

Confusion Matrix for Multiple Classes

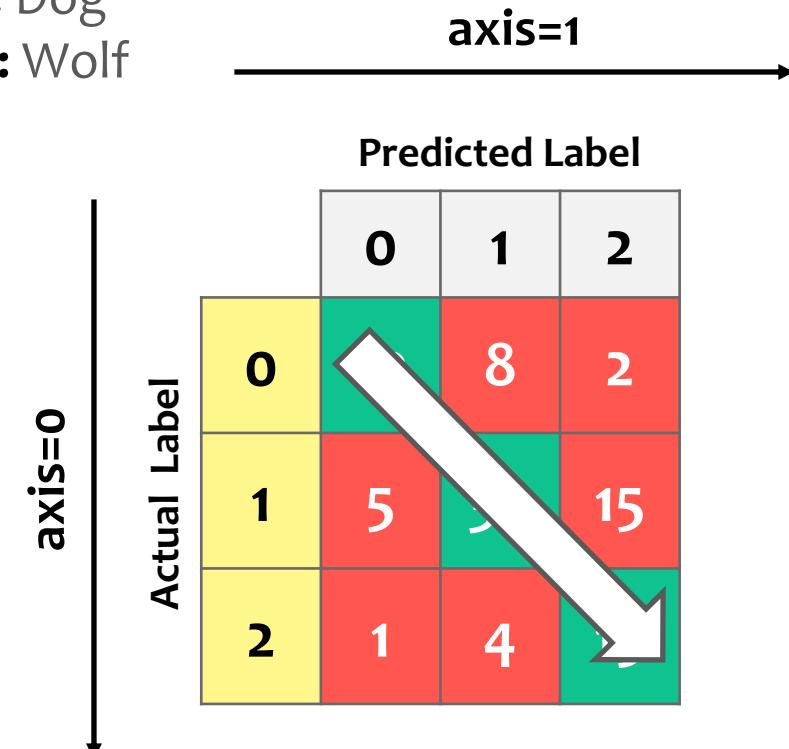
```
from sklearn.metrics import confusion_matrix
import numpy as np

cm = confusion_matrix(y_true, y_pred,
                      labels=['Cat', 'Dog', 'Wolf'])

TP = np.diag(cm)
FP = cm.sum(axis=0) - TP
FN = cm.sum(axis=1) - TP
TN = cm.sum() - (TP + FP + FN)
```

$$TP = \begin{array}{|c|c|c|} \hline 0 & 1 & 2 \\ \hline 10 & 30 & 15 \\ \hline \end{array}$$

0: Cat
1: Dog
2: Wolf



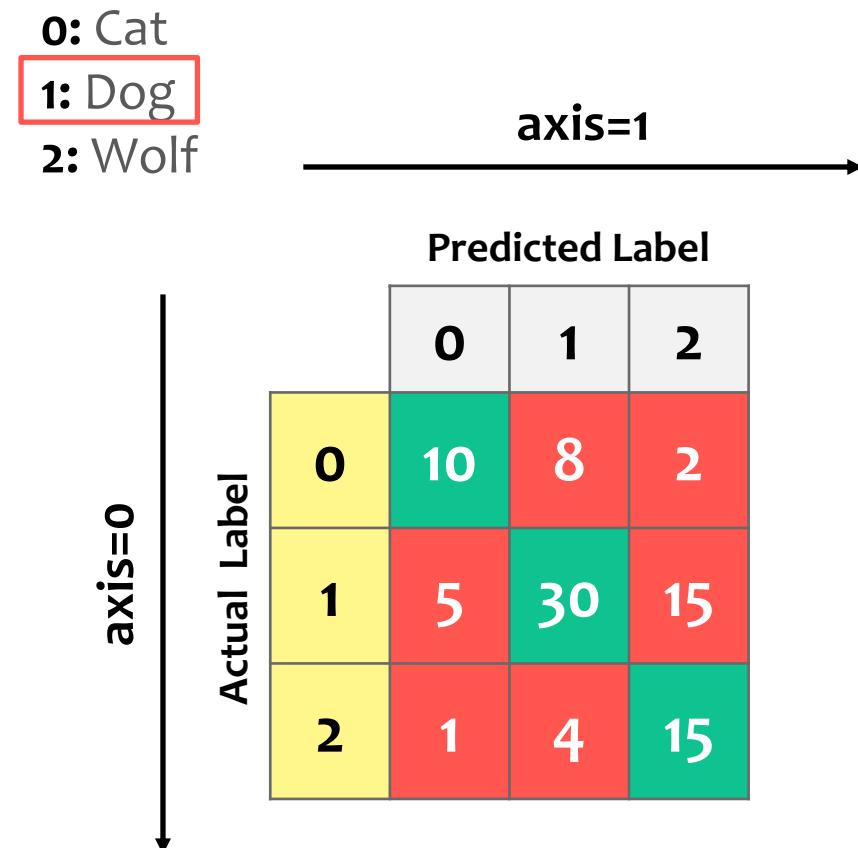
Confusion Matrix for Multiple Classes

```
from sklearn.metrics import confusion_matrix
import numpy as np

cm = confusion_matrix(y_true, y_pred,
                      labels=['Cat', 'Dog', 'Wolf'])

TP = np.diag(cm)
FP = cm.sum(axis=0) - TP
FN = cm.sum(axis=1) - TP
TN = cm.sum() - (TP + FP + FN)
```

$$TP = \begin{array}{ccc} 0 & 1 & 2 \\ \boxed{10} & \boxed{30} & 15 \end{array}$$



$$TP[1] = TP_{Dog}$$

Confusion Matrix for Multiple Classes

```
from sklearn.metrics import confusion_matrix
import numpy as np

cm = confusion_matrix(y_true, y_pred,
                      labels=['Cat', 'Dog', 'Wolf'])

TP = np.diag(cm)
FP = cm.sum(axis=0) - TP
FN = cm.sum(axis=1) - TP
TN = cm.sum() - (TP + FP + FN)
```

$$FP = \begin{matrix} & \begin{matrix} 0 & 1 & 2 \end{matrix} \\ \begin{matrix} 0 & 1 & 2 \end{matrix} & \begin{matrix} 0 & 1 & 2 \end{matrix} \end{matrix} = \begin{matrix} & \begin{matrix} 0 & 1 & 2 \end{matrix} \\ \begin{matrix} 0 & 1 & 2 \end{matrix} & \begin{matrix} 0 & 1 & 2 \end{matrix} \end{matrix}$$

0: Cat
1: Dog
2: Wolf

axis=1

		Predicted Label		
		0	1	2
Actual Label	0	10	8	2
	1	5	30	15
2	1	4	15	

Confusion Matrix for Multiple Classes

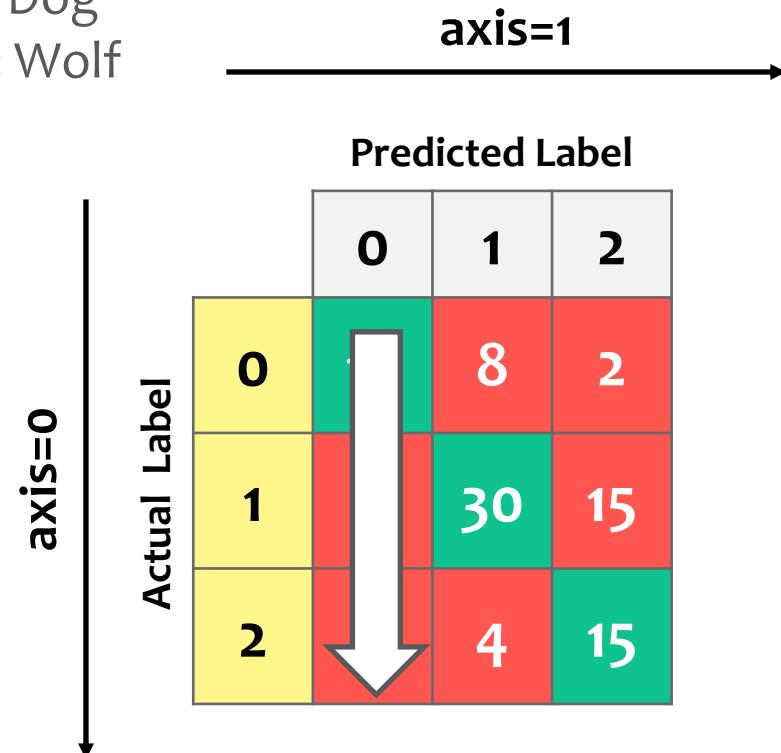
```
from sklearn.metrics import confusion_matrix
import numpy as np

cm = confusion_matrix(y_true, y_pred,
                      labels=['Cat', 'Dog', 'Wolf'])

TP = np.diag(cm)
FP = cm.sum(axis=0) - TP
FN = cm.sum(axis=1) - TP
TN = cm.sum() - (TP + FP + FN)
```

$$FP = \begin{matrix} & \begin{matrix} 0 & 1 & 2 \end{matrix} \\ \begin{matrix} 0 & 1 & 2 \end{matrix} & \begin{matrix} 16 \\ \hline 0 & 1 & 2 \end{matrix} \end{matrix} = \begin{matrix} & \begin{matrix} 0 & 1 & 2 \end{matrix} \\ \begin{matrix} 0 & 1 & 2 \end{matrix} & \begin{matrix} \hline & & \end{matrix} \end{matrix}$$

0: Cat
1: Dog
2: Wolf



Confusion Matrix for Multiple Classes

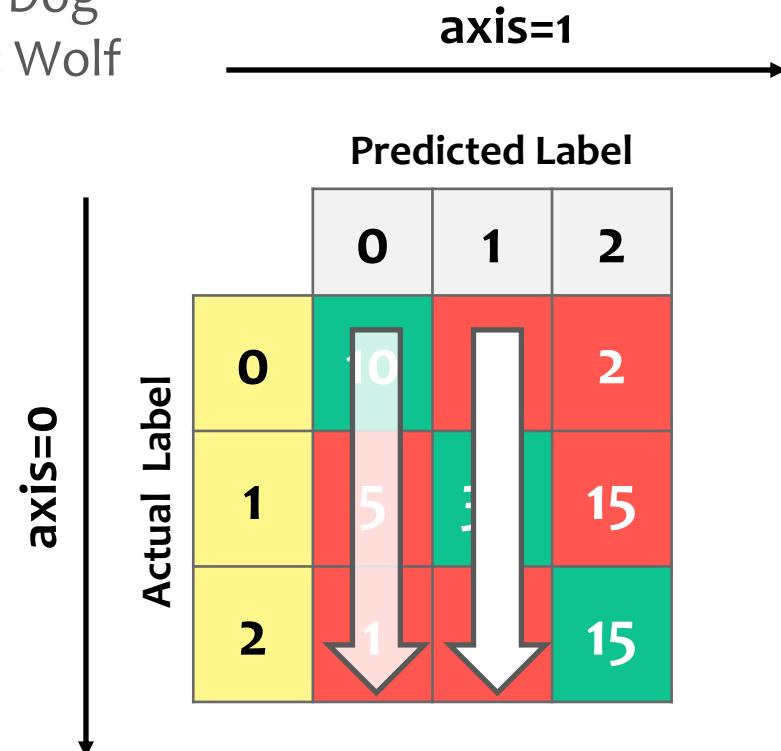
```
from sklearn.metrics import confusion_matrix
import numpy as np

cm = confusion_matrix(y_true, y_pred,
                      labels=['Cat', 'Dog', 'Wolf'])

TP = np.diag(cm)
FP = cm.sum(axis=0) - TP
FN = cm.sum(axis=1) - TP
TN = cm.sum() - (TP + FP + FN)
```

$$FP = \begin{matrix} & \begin{matrix} 0 & 1 & 2 \end{matrix} \\ \begin{matrix} 0 & 1 & 2 \end{matrix} & \begin{matrix} 16 & 42 & \end{matrix} \end{matrix} = \begin{matrix} & \begin{matrix} 0 & 1 & 2 \end{matrix} \\ \begin{matrix} 0 & 1 & 2 \end{matrix} & \begin{matrix} \quad & \quad & \quad \end{matrix} \end{matrix}$$

0: Cat
1: Dog
2: Wolf



Confusion Matrix for Multiple Classes

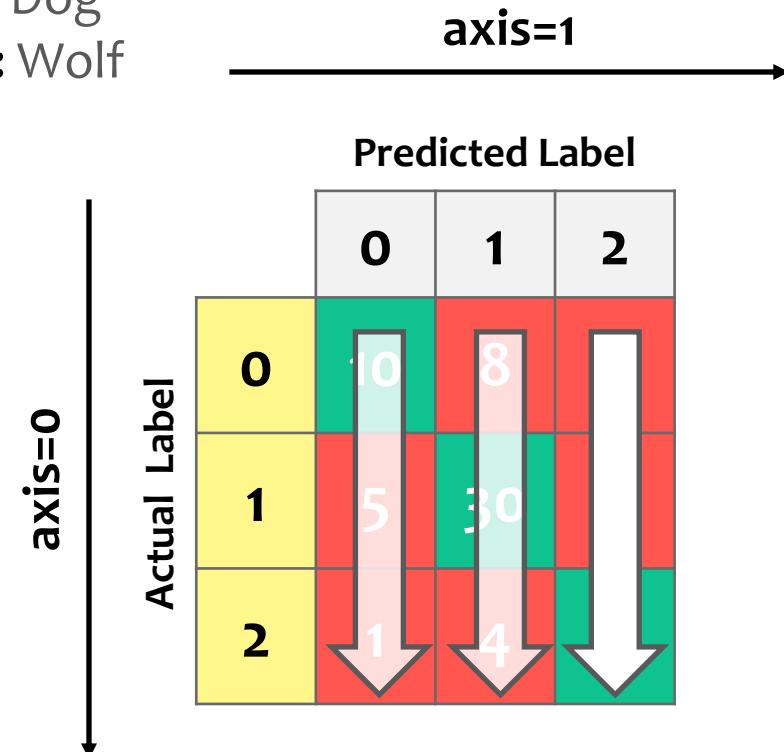
```
from sklearn.metrics import confusion_matrix
import numpy as np

cm = confusion_matrix(y_true, y_pred,
                      labels=['Cat', 'Dog', 'Wolf'])

TP = np.diag(cm)
FP = cm.sum(axis=0) - TP
FN = cm.sum(axis=1) - TP
TN = cm.sum() - (TP + FP + FN)
```

$$FP = \begin{matrix} & \begin{matrix} 0 & 1 & 2 \end{matrix} \\ \begin{matrix} 0 & 1 & 2 \end{matrix} & \begin{matrix} 16 & 42 & 32 \end{matrix} \end{matrix} = \begin{matrix} & \begin{matrix} 0 & 1 & 2 \end{matrix} \\ \begin{matrix} 0 & 1 & 2 \end{matrix} & \begin{matrix} \boxed{} & \boxed{} & \boxed{} \end{matrix} \end{matrix}$$

0: Cat
1: Dog
2: Wolf



Confusion Matrix for Multiple Classes

```
from sklearn.metrics import confusion_matrix
import numpy as np

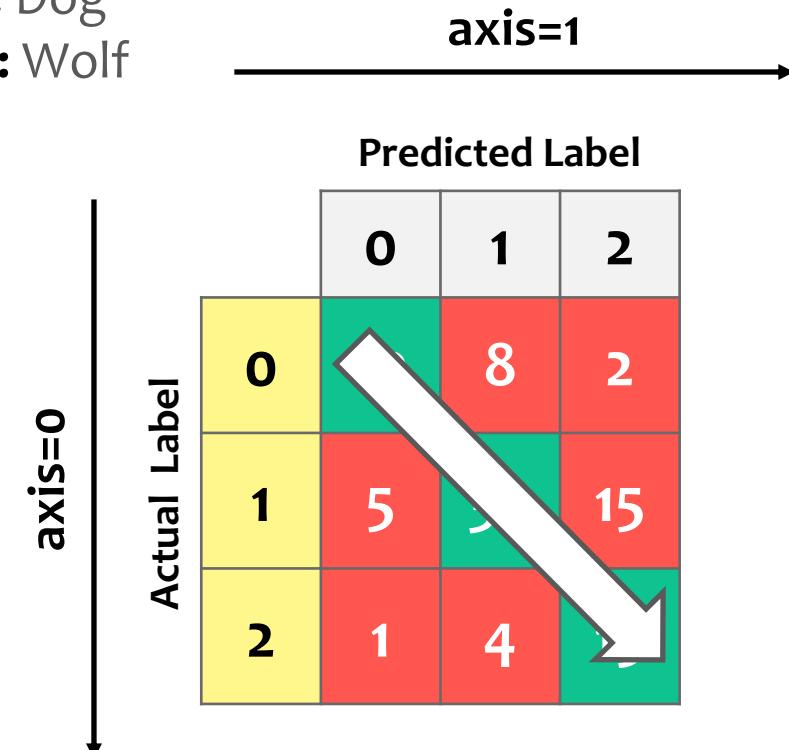
cm = confusion_matrix(y_true, y_pred,
                      labels=['Cat', 'Dog', 'Wolf'])

TP = np.diag(cm)
FP = cm.sum(axis=0) - TP
FN = cm.sum(axis=1) - TP
TN = cm.sum() - (TP + FP + FN)
```

$$FP = \begin{matrix} & \begin{matrix} 0 & 1 & 2 \end{matrix} \\ \begin{matrix} 0 & 1 & 2 \end{matrix} & \begin{matrix} 16 & 42 & 32 \end{matrix} \end{matrix} = \begin{matrix} & \begin{matrix} 0 & 1 & 2 \end{matrix} \\ \begin{matrix} 0 & 1 & 2 \end{matrix} & \begin{matrix} \quad & \quad & \quad \end{matrix} \end{matrix}$$

$$\begin{matrix} & \begin{matrix} 0 & 1 & 2 \end{matrix} \\ \begin{matrix} 0 & 1 & 2 \end{matrix} & \begin{matrix} 10 & 30 & 15 \end{matrix} \end{matrix}$$

0: Cat
1: Dog
2: Wolf



Confusion Matrix for Multiple Classes

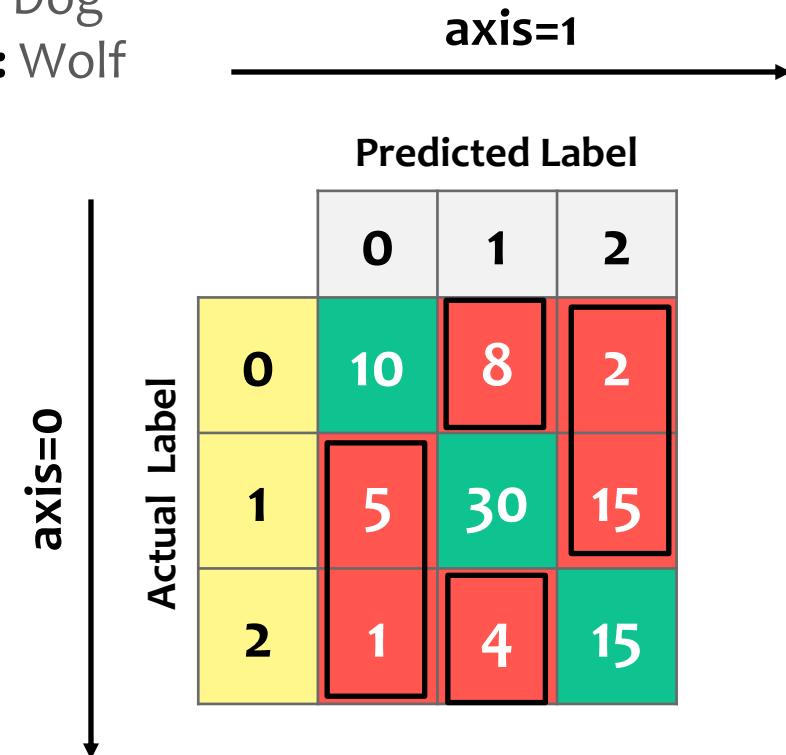
```
from sklearn.metrics import confusion_matrix
import numpy as np

cm = confusion_matrix(y_true, y_pred,
                      labels=['Cat', 'Dog', 'Wolf'])

TP = np.diag(cm)
FP = cm.sum(axis=0) - TP
FN = cm.sum(axis=1) - TP
TN = cm.sum() - (TP + FP + FN)
```

$$FP = \begin{matrix} & \begin{matrix} 0 & 1 & 2 \end{matrix} \\ \begin{matrix} 0 & 10 & 30 & 15 \end{matrix} & \begin{matrix} 16 & 42 & 32 \end{matrix} \end{matrix} = \begin{matrix} & \begin{matrix} 0 & 1 & 2 \end{matrix} \\ \begin{matrix} 0 & 1 & 2 \end{matrix} & \begin{matrix} 6 & 12 & 17 \end{matrix} \end{matrix}$$

0: Cat
1: Dog
2: Wolf



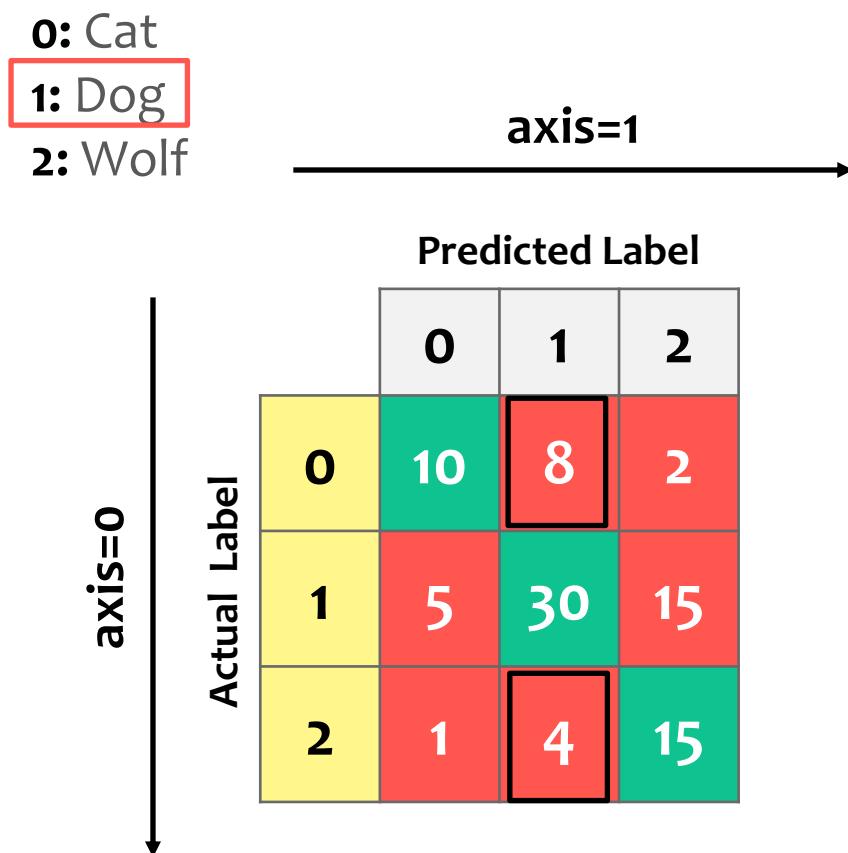
Confusion Matrix for Multiple Classes

```
from sklearn.metrics import confusion_matrix
import numpy as np

cm = confusion_matrix(y_true, y_pred,
                      labels=['Cat', 'Dog', 'Wolf'])

TP = np.diag(cm)
FP = cm.sum(axis=0) - TP
FN = cm.sum(axis=1) - TP
TN = cm.sum() - (TP + FP + FN)
```

$$FP = \frac{\begin{array}{ccc} 0 & 1 & 2 \\ \hline 16 & 42 & 32 \end{array}}{\begin{array}{ccc} 0 & 1 & 2 \\ \hline 10 & 30 & 15 \end{array}} = \begin{array}{ccc} 0 & 1 & 2 \\ \hline 6 & 12 & 17 \end{array}$$



$$FP[1] = FP_{Dog}$$

Confusion Matrix for Multiple Classes

```
from sklearn.metrics import confusion_matrix
import numpy as np

cm = confusion_matrix(y_true, y_pred,
                      labels=['Cat', 'Dog', 'Wolf'])

TP = np.diag(cm)
FP = cm.sum(axis=0) - TP
FN = cm.sum(axis=1) - TP
TN = cm.sum() - (TP + FP + FN)
```

$$FN = \frac{\text{Actual Label } 1}{\text{Predicted Label } 1} = \frac{\text{Actual Label } 2}{\text{Predicted Label } 2}$$

0 1 2

0 1 2

0 1 2

0: Cat
1: Dog
2: Wolf

axis=1

Predicted Label

		0	1	2
Actual Label	0	10	8	2
	1	5	30	15
2	1	4	15	

Confusion Matrix for Multiple Classes

```
from sklearn.metrics import confusion_matrix
import numpy as np

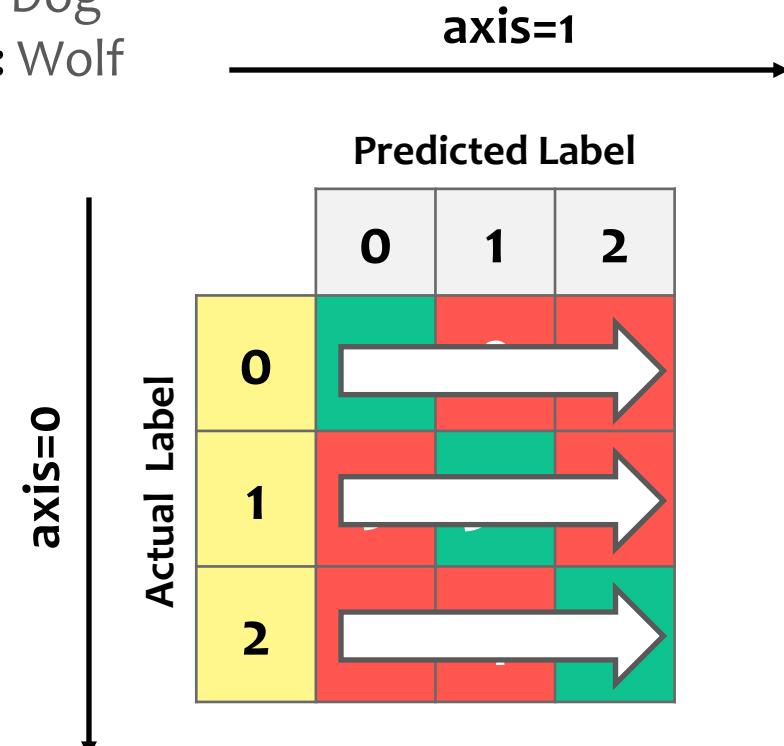
cm = confusion_matrix(y_true, y_pred,
                      labels=['Cat', 'Dog', 'Wolf'])

TP = np.diag(cm)
FP = cm.sum(axis=0) - TP
FN = cm.sum(axis=1) - TP
TN = cm.sum() - (TP + FP + FN)
```

0	1	2
20	50	20

$$FN = \frac{50}{\frac{20+50+20}{3}} = \frac{50}{\frac{90}{3}} = \frac{50}{30} = \frac{5}{3}$$

0: Cat
1: Dog
2: Wolf



Confusion Matrix for Multiple Classes

```
from sklearn.metrics import confusion_matrix
import numpy as np

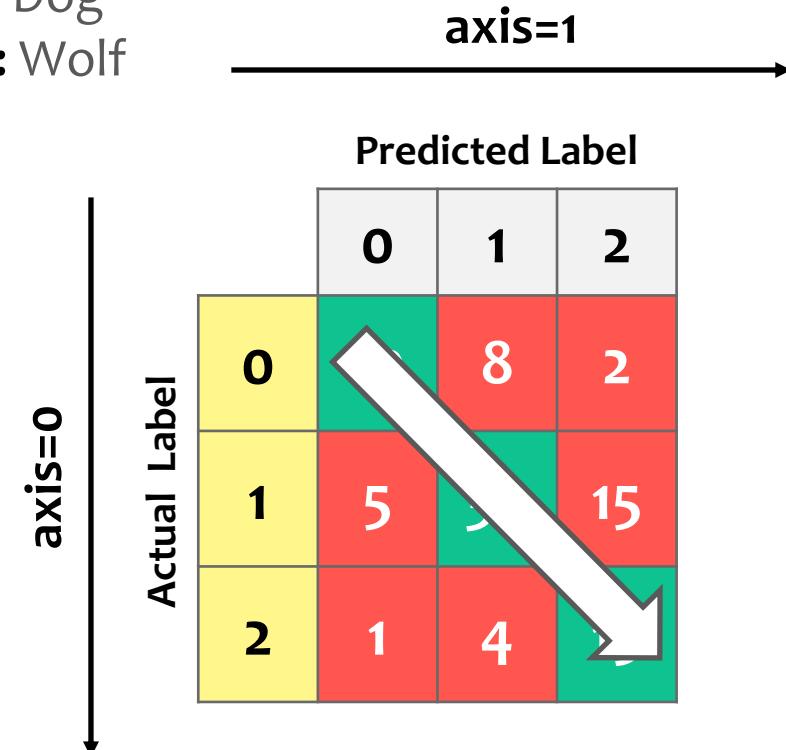
cm = confusion_matrix(y_true, y_pred,
                      labels=['Cat', 'Dog', 'Wolf'])

TP = np.diag(cm)
FP = cm.sum(axis=0) - TP
FN = cm.sum(axis=1) - TP
TN = cm.sum() - (TP + FP + FN)
```

0	1	2
20	50	20

$$FN = \frac{10}{10 + 30 + 15} = \frac{10}{65} = 0.154$$

0: Cat
1: Dog
2: Wolf



Confusion Matrix for Multiple Classes

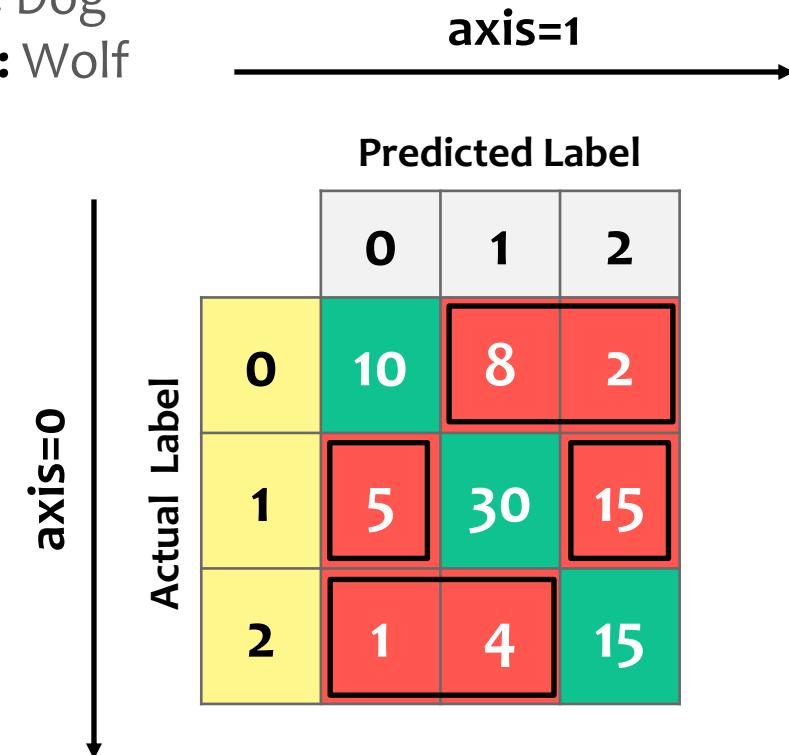
```
from sklearn.metrics import confusion_matrix
import numpy as np

cm = confusion_matrix(y_true, y_pred,
                      labels=['Cat', 'Dog', 'Wolf'])

TP = np.diag(cm)
FP = cm.sum(axis=0) - TP
FN = cm.sum(axis=1) - TP
TN = cm.sum() - (TP + FP + FN)
```

$$FN = \frac{\begin{array}{|c|c|c|} \hline & 0 & 1 & 2 \\ \hline 0 & 20 & 50 & 20 \\ \hline 1 & 10 & 30 & 15 \\ \hline 2 & 1 & 4 & 15 \\ \hline \end{array}}{\begin{array}{|c|c|c|} \hline & 0 & 1 & 2 \\ \hline 0 & 10 & 20 & 5 \\ \hline 1 & 5 & 30 & 15 \\ \hline 2 & 1 & 4 & 15 \\ \hline \end{array}} = \begin{array}{|c|c|c|} \hline & 0 & 1 & 2 \\ \hline 0 & 10 & 20 & 5 \\ \hline 1 & 5 & 30 & 15 \\ \hline 2 & 1 & 4 & 15 \\ \hline \end{array}$$

0: Cat
1: Dog
2: Wolf



Confusion Matrix for Multiple Classes

```
from sklearn.metrics import confusion_matrix
import numpy as np

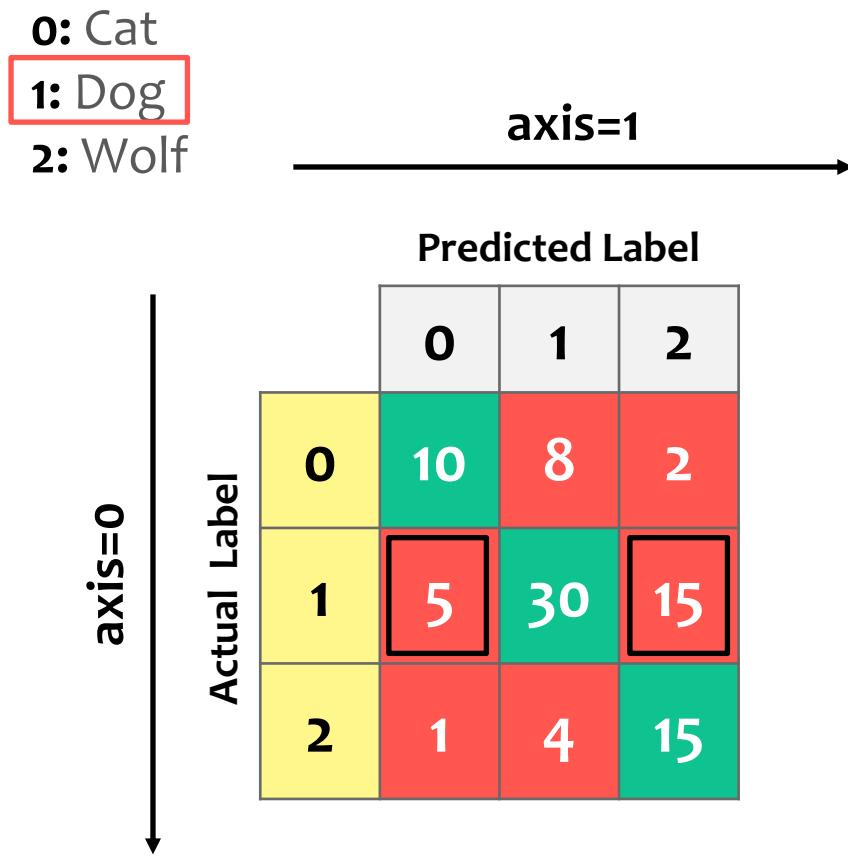
cm = confusion_matrix(y_true, y_pred,
                      labels=['Cat', 'Dog', 'Wolf'])

TP = np.diag(cm)
FP = cm.sum(axis=0) - TP
FN = cm.sum(axis=1) - TP
TN = cm.sum() - (TP + FP + FN)
```

$$FN = \frac{10}{10+30+15} = \frac{10}{55} = 0.18$$

0	1	2
20	50	20

0	1	2
10	20	5



$$FN[1] = FN_{Dog}$$

Confusion Matrix for Multiple Classes

```
from sklearn.metrics import confusion_matrix
import numpy as np

cm = confusion_matrix(y_true, y_pred,
                      labels=['Cat', 'Dog', 'Wolf'])

TP = np.diag(cm)
FP = cm.sum(axis=0) - TP
FN = cm.sum(axis=1) - TP
TN = cm.sum() - (TP + FP + FN)
```

$TN =$

0: Cat
1: Dog
2: Wolf

axis=1

		Predicted Label		
		0	1	2
Actual Label	0	10	8	2
	1	5	30	15
	2	1	4	15

Confusion Matrix for Multiple Classes

```
from sklearn.metrics import confusion_matrix
import numpy as np

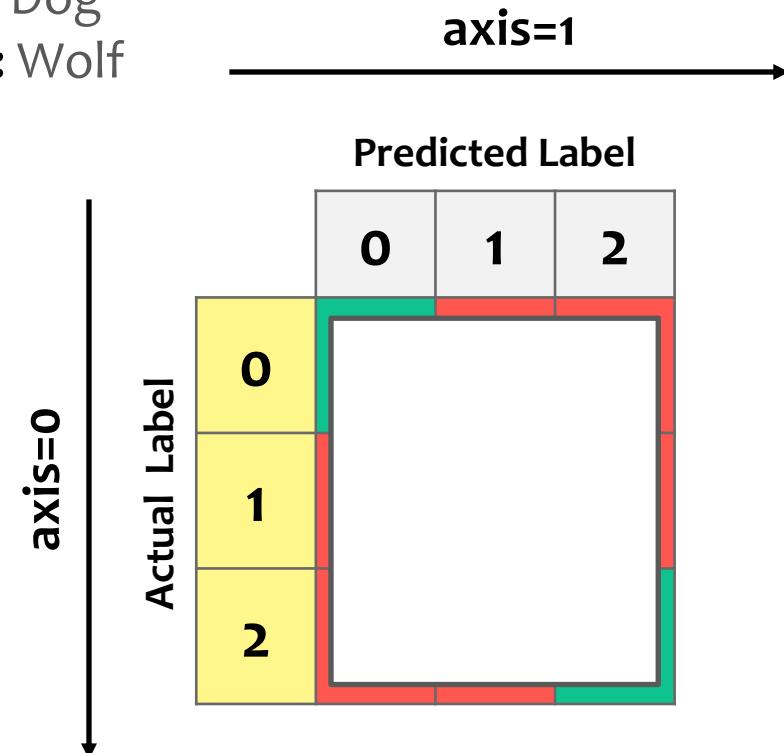
cm = confusion_matrix(y_true, y_pred,
                      labels=['Cat', 'Dog', 'Wolf'])

TP = np.diag(cm)
FP = cm.sum(axis=0) - TP
FN = cm.sum(axis=1) - TP
TN = cm.sum() - (TP + FP + FN)
```

90

$TN =$

0: Cat
1: Dog
2: Wolf



Confusion Matrix for Multiple Classes

```
from sklearn.metrics import confusion_matrix
import numpy as np

cm = confusion_matrix(y_true, y_pred,
                      labels=['Cat', 'Dog', 'Wolf'])

TP = np.diag(cm)
FP = cm.sum(axis=0) - TP
FN = cm.sum(axis=1) - TP
TN = cm.sum() - (TP + FP + FN)
```

$$TN = \begin{bmatrix} & \begin{matrix} 0 & 1 & 2 \end{matrix} \\ \begin{matrix} 90 \\ - \end{matrix} & \begin{bmatrix} TP & \begin{matrix} 10 & 30 & 15 \end{matrix} \\ + & \begin{bmatrix} 6 & 12 & 17 \end{bmatrix} \\ FP & \begin{bmatrix} 10 & 20 & 5 \end{bmatrix} \\ FN & \end{bmatrix} \end{bmatrix}$$

0: Cat
1: Dog
2: Wolf

axis=1

		Predicted Label		
		0	1	2
Actual Label	0	10	8	2
	1	5	30	15
2	17	4	15	

Confusion Matrix for Multiple Classes

```
from sklearn.metrics import confusion_matrix
import numpy as np

cm = confusion_matrix(y_true, y_pred,
                      labels=['Cat', 'Dog', 'Wolf'])

TP = np.diag(cm)
FP = cm.sum(axis=0) - TP
FN = cm.sum(axis=1) - TP
TN = cm.sum() - (TP + FP + FN)
```

$$TN = \begin{matrix} & \begin{matrix} 90 \end{matrix} \\ \begin{matrix} 0 & 1 & 2 \end{matrix} & \begin{matrix} 26 & 62 & 37 \end{matrix} \end{matrix} = \begin{matrix} & \begin{matrix} 0 & 1 & 2 \end{matrix} \\ \begin{matrix} 0 & 1 & 2 \end{matrix} & \boxed{\quad \quad \quad} \end{matrix}$$

0: Cat
1: Dog
2: Wolf

axis=1

		Predicted Label		
		0	1	2
Actual Label	0	10	8	2
	1	5	30	15
2	2	1	4	15

Confusion Matrix for Multiple Classes

```
from sklearn.metrics import confusion_matrix
import numpy as np

cm = confusion_matrix(y_true, y_pred,
                      labels=['Cat', 'Dog', 'Wolf'])

TP = np.diag(cm)
FP = cm.sum(axis=0) - TP
FN = cm.sum(axis=1) - TP
TN = cm.sum() - (TP + FP + FN)
```

$$TN = \begin{matrix} & \begin{matrix} 90 \end{matrix} \\ \begin{matrix} 0 & 1 & 2 \end{matrix} & \begin{matrix} \boxed{26} & \boxed{62} & \boxed{37} \end{matrix} \end{matrix} = \begin{matrix} & \begin{matrix} 0 & 1 & 2 \end{matrix} \\ \begin{matrix} 0 & 1 & 2 \end{matrix} & \begin{matrix} \boxed{64} & \boxed{28} & \boxed{53} \end{matrix} \end{matrix}$$

0: Cat
1: Dog
2: Wolf

axis=1

		Predicted Label		
		0	1	2
Actual Label	0	10	8	2
	1	5	30	15
2	1	4	15	

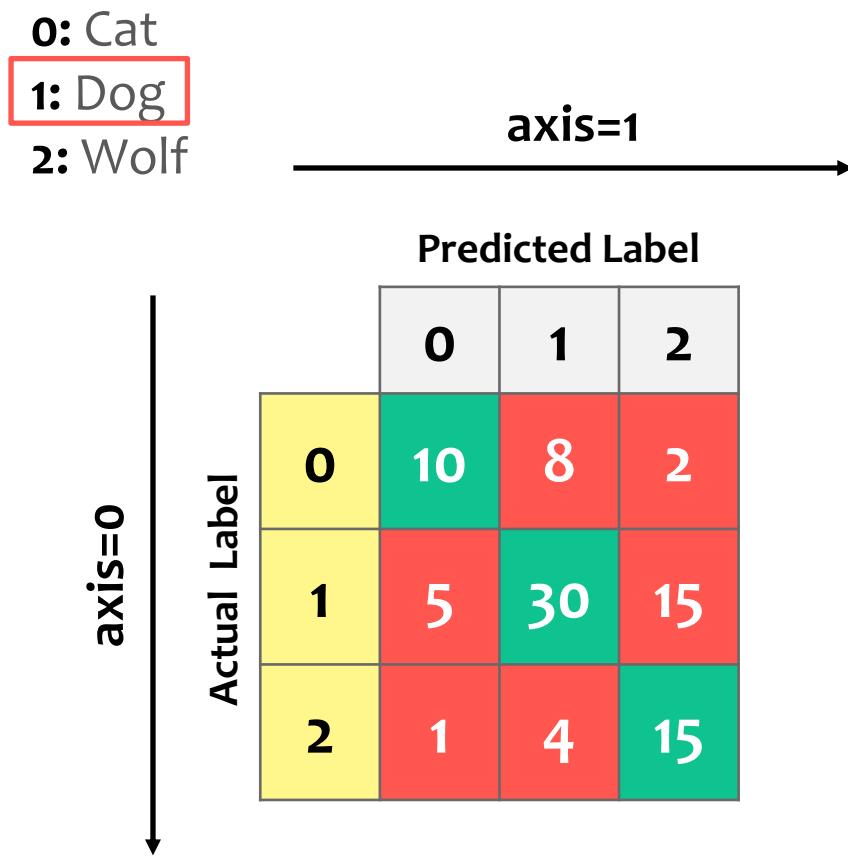
Confusion Matrix for Multiple Classes

```
from sklearn.metrics import confusion_matrix
import numpy as np

cm = confusion_matrix(y_true, y_pred,
                      labels=['Cat', 'Dog', 'Wolf'])

TP = np.diag(cm)
FP = cm.sum(axis=0) - TP
FN = cm.sum(axis=1) - TP
TN = cm.sum() - (TP + FP + FN)
```

$$TN = \begin{matrix} & \begin{matrix} 90 \end{matrix} \\ \begin{matrix} 0 & 1 & 2 \end{matrix} & \begin{matrix} \boxed{26} & 62 & 37 \end{matrix} \end{matrix} = \begin{matrix} & \begin{matrix} 0 & 1 & 2 \end{matrix} \\ \begin{matrix} 0 & \boxed{1} & 2 \end{matrix} & \begin{matrix} 64 & \boxed{28} & 53 \end{matrix} \end{matrix}$$



$$TN[1] = TN_{Dog}$$

Sensitivity vs Specificity

		Predicted Label	
		Negative	Positive
Actual Label	Healthy	Healthy	Cancer
	Healthy	TN	FP
	Cancer	FN	TP

Sensitivity (Recall or True Positive Rate)

What is the **proportion** of **positives** that were **correctly classified** as **positives**?

What is the **proportion** of **patients with cancer** that were **correctly identified**?

How sensitive is the classifier **to correctly identify patients with cancer**?

$$Sensitivity = Recall = \frac{TP}{TP + FN}$$

		Predicted Label	
		Negative	Positive
Actual Label	Healthy	Healthy	Cancer
	Healthy	TN	FP
	Cancer	FN	TP

Sensitivity (Recall or True Positive Rate)

What is the **proportion** of **positives** that were **correctly classified** as **positives**?

What is the **proportion** of **patients with cancer** that were **correctly identified**?

How sensitive is the classifier **to correctly identify patients with cancer**?

$$Sensitivity = Recall = \frac{TP}{TP + FN}$$

⚠ It can cheat by guessing everything as positive (recall=1).

		Predicted Label	
		Negative	Positive
Actual Label	Healthy	Cancer	
	Healthy	TN	FP
Cancer	FN		TP

Sensitivity (Recall or True Positive Rate)

What is the **proportion** of **positives** that were **correctly classified** as **positives**?

What is the **proportion** of **patients with cancer** that were **correctly identified**?

How sensitive is the classifier **to correctly identify patients with cancer**?

$$Sensitivity = Recall = \frac{TP}{TP + FN}$$

`sklearn.metrics.recall_score`

```
from sklearn.metrics import recall_score  
recall_score(y_true, y_pred)
```



By default, the **positive class** must have **value/label +1**.

Sensitivity (Recall or True Positive Rate)

What is the **proportion** of **positives** that were **correctly classified** as **positives**?

What is the **proportion** of **patients with cancer** that were **correctly identified**?

How sensitive is the classifier **to correctly identify patients with cancer**?

$$Sensitivity = Recall = \frac{TP}{TP + FN}$$

sklearn.metrics.recall_score

```
from sklearn.metrics import recall_score  
recall_score(y_true, y_pred)
```

```
from sklearn.metrics import recall_score  
recall_score(y_true, y_pred, pos_label='Cancer')
```



By default, the **positive class** must have **value/label +1**.



We can **indicate** which is the **value/label** of the **positive class**.

Specificity (True Negative Rate)

What is the **proportion** of **negatives** that were **correctly classified** as **negatives**?

What is the **proportion** of **healthy patients** that were **correctly identified**?

$$\text{Specificity} = \frac{\text{TN}}{\text{TN} + \text{FP}}$$

		Predicted Label	
		Negative	Positive
Actual Label	Healthy	Healthy	Cancer
	Healthy	TN	FP
	Cancer	FN	TP

Specificity (True Negative Rate)

What is the **proportion** of **negatives** that were **correctly classified** as **negatives**?

What is the **proportion** of **healthy patients** that were **correctly identified**?

$$\text{Specificity} = \frac{\text{TN}}{\text{TN} + \text{FP}}$$

		Predicted Label	
		Negative	Positive
Actual Label Positive	Healthy	Healthy	Cancer
	Healthy	TN	FP
	Cancer	FN	TP

 It can **cheat** by guessing **everything** as **negative** (specificity=1).

Specificity (True Negative Rate)

What is the **proportion** of **negatives** that were **correctly classified** as **negatives**?

What is the **proportion** of **healthy patients** that were **correctly identified**?

$$\text{Specificity} = \frac{TN}{TN + FP}$$

```
from sklearn.metrics import confusion_matrix  
  
tn, fp, fn, tp = confusion_matrix(y_true, y_pred).ravel()  
  
specificity = tn / (tn + fp)
```

```
from sklearn.metrics import confusion_matrix  
  
tn, fp, fn, tp = confusion_matrix(y_true, y_pred,  
                                    labels=['Healthy', 'Cancer']).ravel()  
  
specificity = tn / (tn + fp)
```

Specificity (True Negative Rate)

What is the **proportion** of **negatives** that were **correctly classified** as **negatives**?

What is the **proportion** of **healthy patients** that were **correctly identified**?

$$\text{Specificity} = \frac{TN}{TN + FP}$$



For the **binary case**, we could say that **specificity** is the '**sensitivity/recall**' of the **negative class**.



If we replace N with P, we would have:

$$\frac{TP}{TP + FN}$$

which is exactly the formula for '**sensitivity/recall**'.

Specificity (True Negative Rate)

What is the **proportion** of **negatives** that were **correctly classified** as **negatives**?

What is the **proportion** of **healthy patients** that were **correctly identified**?

$$\text{Specificity} = \frac{TN}{TN + FP}$$



For the **binary case**, we could say that **specificity** is the '**sensitivity/recall**' of the **negative class**.



If we replace N with P, we would have:

$$\frac{TP}{TP + FN}$$

which is exactly the formula for '**sensitivity/recall**'.

sklearn.metrics.recall_score

```
from sklearn.metrics import recall_score  
  
# for two classes, we can compute sensitivity by assuming  
# the negative class as the positive class in `recall_score`.  
recall_score(y_true_bin, y_pred_bin_knn, pos_label='Healthy')
```

		Predicted Label	
		Negative	Positive
Actual Label	Healthy	Healthy	Cancer
	Healthy	70	130
		Cancer	20
			80

$$Sensitivity = \frac{TP}{TP + FN}$$

		Predicted Label	
		Negative	Positive
Actual Label	Healthy	Healthy	Cancer
	Healthy	70	130
		Cancer	80
		20	

$$Sensitivity = \frac{80}{80 + 20} = 0.8$$

80% of the people with cancer
were correctly identified.

		Predicted Label	
		Negative	Positive
		Healthy	Cancer
Actual Label	Positive	70	130
	Negative	20	80

$$Sensitivity = \frac{80}{80 + 20} = 0.8$$

$$Specificity = \frac{TN}{TN + FP}$$

		Predicted Label	
		Negative	Positive
Actual Label	Healthy	Healthy	Cancer
	Healthy	70	130
Positive	Cancer	20	80

$$Sensitivity = \frac{80}{80 + 20} = 0.8$$

$$Specificity = \frac{70}{70 + 130} = 0.35$$

35% of the healthy people were correctly identified.

		Predicted Label	
		Negative	Positive
Actual Label	Healthy	Healthy	Cancer
	Healthy	70	130
Positive		20	80

Confusion matrix for classifier A
(e.g., KNN)

		Predicted Label	
		Negative	Positive
		Healthy	Cancer
Actual Label Positive	Healthy	70	130
	Cancer	20	80

Confusion matrix for classifier B
(e.g., Logistic Regression)

		Predicted Label	
		Negative	Positive
		Healthy	Cancer
Actual Label Positive	Healthy	80	120
	Cancer	50	50

Confusion matrix for classifier A
(e.g., KNN)

		Predicted Label	
		Negative	Positive
		Healthy	Cancer
Actual Label Positive	Healthy	70	130
	Cancer	20	80

$$Sensitivity = \frac{80}{80 + 20} = 0.8$$

$$Specificity = \frac{70}{70 + 130} = 0.35$$

Confusion matrix for classifier B
(e.g., Logistic Regression)

		Predicted Label	
		Negative	Positive
		Healthy	Cancer
Actual Label Positive	Healthy	80	120
	Cancer	50	50

$$Sensitivity = \frac{50}{50 + 50} = 0.5$$

$$Specificity = \frac{80}{80 + 120} = 0.4$$

Confusion matrix for classifier A
(e.g., KNN)

		Predicted Label	
		Negative	Positive
Actual Label	Healthy	Healthy	Cancer
	Cancer	20	80
Positive	70	130	
Negative			130

$$Sensitivity = \frac{80}{80 + 20} = 0.8$$

Sensitivity tells us that KNN is better at **correctly classifying positives (people with cancer)**

Confusion matrix for classifier B
(e.g., Logistic Regression)

		Predicted Label	
		Negative	Positive
Actual Label	Healthy	Healthy	Cancer
	Cancer	50	50
Positive	80	120	
Negative			120

$$Sensitivity = \frac{50}{50 + 50} = 0.5$$

$$Specificity = \frac{80}{80 + 120} = 0.4$$

Confusion matrix for classifier A
(e.g., KNN)

		Predicted Label	
		Negative	Positive
		Healthy	Cancer
Actual Label Positive	Healthy	70	130
	Cancer	20	80

$$Sensitivity = \frac{80}{80 + 20} = 0.8$$

$$Specificity = \frac{70}{70 + 130} = 0.35$$

Confusion matrix for classifier B
(e.g., Logistic Regression)

		Predicted Label	
		Negative	Positive
		Healthy	Cancer
Actual Label Positive	Healthy	80	120
	Cancer	50	50

Specificity tells us that **Logistic Regression** is better at **correctly classifying negatives (healthy people)**

$$Specificity = \frac{80}{80 + 120} = 0.4$$

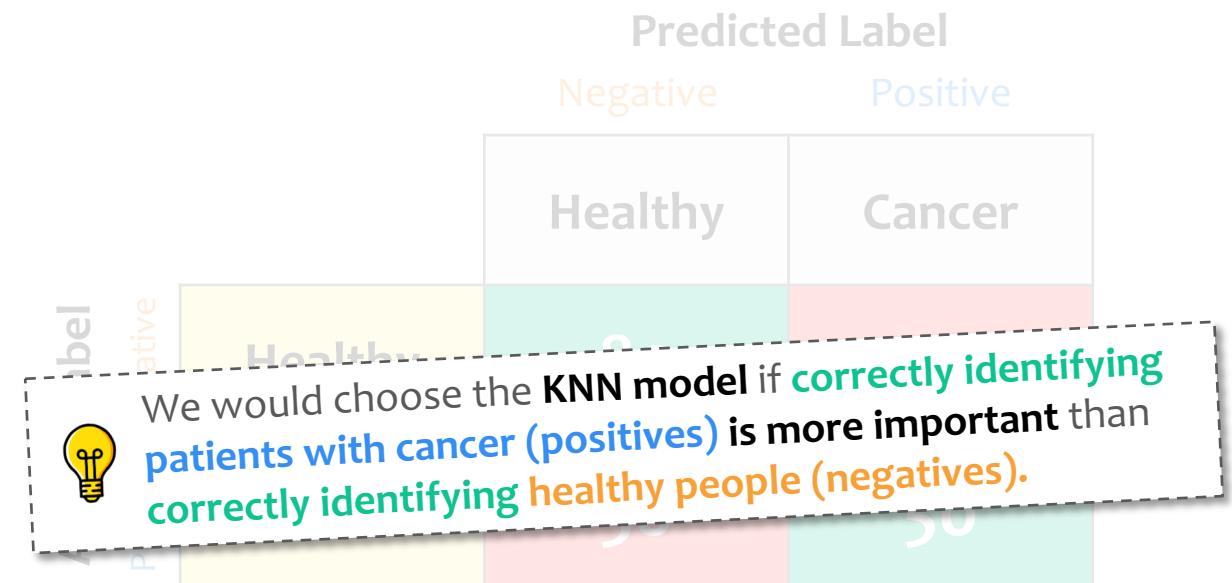
Confusion matrix for classifier A
(e.g., KNN)

		Predicted Label	
		Negative	Positive
Actual Label	Healthy	Healthy	Cancer
	Cancer	20	80

$$Sensitivity = \frac{80}{80 + 20} = 0.8$$

$$Specificity = \frac{70}{70 + 130} = 0.35$$

Confusion matrix for classifier B
(e.g., Logistic Regression)



$$Sensitivity = \frac{50}{50 + 50} = 0.5$$

$$Specificity = \frac{80}{80 + 120} = 0.4$$

Confusion matrix for classifier A (e.g., KNN)

		Predicted Label	
		Negative	Positive
Actual Label	Healthy	80	20
	Cancer	70	130

 Otherwise, we would choose the **Logistic model** if **correctly identifying healthy people (negatives) is more important than correctly identifying patients with cancer (positives).**

$$Sensitivity = \frac{80}{80 + 20} = 0.8$$

$$Specificity = \frac{70}{70 + 130} = 0.35$$

Confusion matrix for classifier B (e.g., Logistic Regression)

		Predicted Label	
		Negative	Positive
Actual Label	Healthy	80	120
	Cancer	50	50

$$Sensitivity = \frac{50}{50 + 50} = 0.5$$

$$Specificity = \frac{80}{80 + 120} = 0.4$$

Sensitivity for Multiple Classes

		Predicted Label		
		Cat	Dog	Wolf
Actual Label	Cat	TP _{Cat}		
	Dog		TP _{Dog}	
	Wolf			TP _{Wolf}

Sensitivity for Multiple Classes

		Predicted Label		
		Cat	Dog	Wolf
Actual Label	Cat	10	8	2
	Dog	5	30	15
	Wolf	1	4	15

Sensitivity for Multiple Classes

$$Sensitivity_{Cat} = Recall_{Cat} = \frac{TP_{Cat}}{TP_{Cat} + FN_{Cat}} =$$

		Predicted Label		
		Cat	Dog	Wolf
Actual Label	Cat	10	8	2
	Dog	5	30	15
	Wolf	1	4	15

Sensitivity for Multiple Classes

$$Sensitivity_{Cat} = Recall_{Cat} = \frac{TP_{Cat}}{TP_{Cat} + FN_{Cat}} = \frac{10}{10 +}$$

		Predicted Label		
		Cat	Dog	Wolf
Actual Label	Cat	10	8	2
	Dog	5	30	15
	Wolf	1	4	15

Sensitivity for Multiple Classes

$$Sensitivity_{Cat} = Recall_{Cat} = \frac{TP_{Cat}}{TP_{Cat} + FN_{Cat}} = \frac{10}{10 + 10}$$

		Predicted Label		
		Cat	Dog	Wolf
Actual Label	Cat	10	8	2
	Dog	5	30	15
	Wolf	1	4	15

Sensitivity for Multiple Classes

$$Sensitivity_{Cat} = Recall_{Cat} = \frac{TP_{Cat}}{TP_{Cat} + FN_{Cat}} = \frac{10}{10 + 10} = 0.5$$

 $Sensitivity_{Cat}$ tells us that 50% of cats were correctly identified.

		Predicted Label		
		Cat	Dog	Wolf
Actual Label	Cat	10	8	2
	Dog	5	30	15
	Wolf	1	4	15

Sensitivity for Multiple Classes

$$Sensitivity_{Cat} = Recall_{Cat} = \frac{TP_{Cat}}{TP_{Cat} + FN_{Cat}} = \frac{10}{10 + 10} = 0.5$$



Sensitivity_{Cat} tells us that 50% of cats were correctly identified.

```
from sklearn.metrics import confusion_matrix  
recall_score(y_true, y_pred, average=None, labels=['Cat', 'Dog', 'Wolf'])  
  
array([0.5 , 0.6 , 0.75])
```

	Cat	Dog	Wolf
Cat	10	8	2
Dog	5	30	15
Wolf	1	4	15

Sensitivity for Multiple Classes

$$Sensitivity_{Cat} = Recall_{Cat} = \frac{TP_{Cat}}{TP_{Cat} + FN_{Cat}} = \frac{10}{10 + 10} = 0.5$$



Sensitivity_{Cat} tells us that 50% of cats were correctly identified.

```
from sklearn.metrics import confusion_matrix  
recall_score(y_true, y_pred, average=None, labels=['Cat', 'Dog', 'Wolf'])  
  
array([0.5 , 0.6 , 0.75])
```

	Cat	Dog	Wolf
Cat	10	8	2
Dog	5	30	15
Wolf	1	4	15

Sensitivity for Multiple Classes

$$Sensitivity_{Cat} = Recall_{Cat} = \frac{TP_{Cat}}{TP_{Cat} + FN_{Cat}} = \frac{10}{10 + 10} = 0.5$$



Sensitivity_{Cat} tells us that 50% of cats were correctly identified.

```
from sklearn.metrics import confusion_matrix  
recall_score(y_true, y_pred, average=None, labels=['Cat', 'Dog', 'Wolf'])  
  
array([0.5, 0.6, 0.75])
```

	Cat	Dog	Wolf
Cat	10	8	2
Dog	5	30	15
Wolf	1	4	15

Sensitivity for Multiple Classes

$$Sensitivity_{Cat} = Recall_{Cat} = \frac{TP_{Cat}}{TP_{Cat} + FN_{Cat}} = \frac{10}{10 + 10} = 0.5$$



Sensitivity_{Cat} tells us that 50% of cats were correctly identified.

```
from sklearn.metrics import confusion_matrix  
recall
```

We can return a single “global” sensitivity by averaging the metrics according to some strategy.

		Cat	Dog	Wolf
Actual Label	Cat	10	8	2
	Dog	5	30	15
	Wolf	1	4	15

Single Sensitivity for Multiple Classes

macro

Calculate **metrics** for **each label** and **average** them.

For **c** classes, enumerated from **0** to **c-1**

$$Sensitivity_{macro} = Recall_{macro} = \frac{\sum_{j=0}^{c-1} Sensitivity_j}{c}$$

Single Sensitivity for Multiple Classes

macro

Calculate **metrics** for **each label** and **average** them.

For **c** classes, enumerated from **0** to **c-1**

$$Sensitivity_{macro} = Recall_{macro} = \frac{\sum_{j=0}^{c-1} Sensitivity_j}{c}$$



Each class **contributes/weights equally** to the final metric since it **does not** take **class proportion (imbalance)** into account.
Thus, it is appropriate for **class-imbalanced datasets**.

Single Sensitivity for Multiple Classes

macro

Calculate **metrics** for **each label** and **average** them.

For **c** classes, enumerated from **0** to **c-1**

$$Sensitivity_{macro} = Recall_{macro} = \frac{\sum_{j=0}^{c-1} Sensitivity_j}{c}$$

```
from sklearn.metrics import recall_score  
recall_score(y_true, y_pred, average='macro',  
              labels=['Cat', 'Dog', 'Wolf'])
```

0.6166666666666667



Each class **contributes/weights equally** to the final metric since it **does not** take **class proportion (imbalance)** into account. Thus, it is appropriate for **class-imbalanced datasets**.

Single Sensitivity for Multiple Classes

macro

Calculate **metrics** for **each label** and **average** them.

For **c** classes, enumerated from **0** to **c-1**

$$Sensitivity_{macro} = Recall_{macro} = \frac{\sum_{j=0}^{c-1} Sensitivity_j}{c}$$

```
from sklearn.metrics import recall_score  
recall_score(y_true, y_pred, average='macro',  
              labels=['Cat', 'Dog', 'Wolf'])
```

0.6166666666666667



Each class **contributes/weights equally** to the final metric since it **does not** take **class proportion (imbalance)** into account.
Thus, it is appropriate for **class-imbalanced datasets**.

micro

Calculate **metrics globally** by counting the total true positives, false negatives and false positives.

For **c** classes, enumerated from **0** to **c-1**, and **m** samples

$$Sensitivity_{micro} = Recall_{micro} = \frac{\sum_{j=0}^{c-1} TP_j}{m}$$

Single Sensitivity for Multiple Classes

macro

Calculate **metrics** for **each label** and **average** them.

For **c** classes, enumerated from **0** to **c-1**

$$Sensitivity_{macro} = Recall_{macro} = \frac{\sum_{j=0}^{c-1} Sensitivity_j}{c}$$



Each class **contributes/weights equally** to the final metric since it **does not** take **class proportion (imbalance)** into account.
Thus, it is appropriate for **class-imbalanced datasets**.

micro

Calculate **metrics globally** by counting the total true positives, false negatives and false positives.

For **c** classes, enumerated from **0** to **c-1**, and **m** samples

$$Sensitivity_{micro} = Recall_{micro} = \frac{\sum_{j=0}^{c-1} TP_j}{m}$$

```
from sklearn.metrics import recall_score  
recall_score(y_true, y_pred, average='macro',  
              labels=['Cat', 'Dog', 'Wolf'])
```

0.6166666666666667

```
from sklearn.metrics import recall_score  
recall_score(y_true, y_pred, average='micro',  
              labels=['Cat', 'Dog', 'Wolf'])
```

0.6111111111111112

Single Sensitivity for Multiple Classes

weighted

Calculate metrics for each class, and find their **average weighted** by the **number of true samples of each class**.

Let c be the number of classes, enumerated from **0** to **c-1**,
 m the number samples, and m_j is the number of true samples of class j .

$$Sensitivity_{weighted} = Recall_{weighted} = \frac{\sum_{j=0}^{c-1} m_j * Sensitivity_j}{m}$$

Single Sensitivity for Multiple Classes

weighted

Calculate metrics for each class, and find their **average weighted** by the **number of true samples of each class**.

Let c be the number of classes, enumerated from **0** to **c-1**,
 m the number samples, and m_j is the number of true samples of class j .

$$Sensitivity_{weighted} = Recall_{weighted} = \frac{\sum_{j=0}^{c-1} m_j * Sensitivity_j}{m}$$

weight (w_j) for class j

Single Sensitivity for Multiple Classes

weighted

Calculate metrics for each class, and find their **average weighted** by the **number of true samples of each class**.

Let c be the number of classes, enumerated from **0** to **$c-1$** ,
 m the number samples, and m_j is the number of true samples of class j .

$$Sensitivity_{weighted} = Recall_{weighted} = \frac{\sum_{j=0}^{c-1} m_j * Sensitivity_j}{m}$$

weight (w_j) for class j



It take into account class proportion (imbalance) for computation.
Thus, it is **NOT appropriated** for class-imbalanced datasets.

Single Sensitivity for Multiple Classes

weighted

Calculate metrics for each class, and find their **average weighted** by the **number of true samples of each class**.

Let c be the number of classes, enumerated from **0** to **$c-1$** ,
 m the number samples, and m_j is the number of true samples of class j .

```
from sklearn.metrics import recall_score  
recall_score(y_true, y_pred, average='weighted',  
              labels=['Cat', 'Dog', 'Wolf'])
```

0.611111111111112

$$Sensitivity_{weighted} = Recall_{weighted} = \frac{\sum_{j=0}^{c-1} m_j * Sensitivity_j}{m}$$

weight (w_j) for class j



It take into account class proportion (imbalance) for computation.
Thus, it is **NOT appropriated** for class-imbalanced datasets.

Single Sensitivity for Multiple Classes

weighted

Calculate metrics for each class, and find their **average weighted** by the **number of true samples of each class**.

Let c be the number of classes, enumerated from **0** to **$c-1$** ,
 m the number samples, and m_j is the number of true samples of class j .

$$Sensitivity_{\text{weighted}} = Recall_{\text{weighted}} = \frac{\sum_{j=0}^{c-1} m_j * Sensitivity_j}{m}$$

weight (w_j) for class j



It take into account class proportion (imbalance) for computation.
Thus, it is **NOT appropriated** for class-imbalanced datasets.



$Sensitivity_{\text{weighted}}$ is equal to **accuracy**.

```
from sklearn.metrics import recall_score  
recall_score(y_true, y_pred, average='weighted',  
              labels=['Cat', 'Dog', 'Wolf'])
```

0.6111111111111112

Specificity for Multiple Classes

$$Specificity_{Cat} = \frac{TN_{Cat}}{TN_{Cat} + FP_{Cat}}$$

		Predicted Label		
		Cat	Dog	Wolf
Actual Label	Cat	10	8	2
	Dog	5	30	15
	Wolf	1	4	15

Specificity for Multiple Classes

$$Specificity_{Cat} = \frac{TN_{Cat}}{TN_{Cat} + FP_{Cat}} = \frac{64}{64 + }$$

		Predicted Label		
		Cat	Dog	Wolf
Actual Label	Cat	10	8	2
	Dog	5	30	15
	Wolf	1	4	15

Specificity for Multiple Classes

$$Specificity_{Cat} = \frac{TN_{Cat}}{TN_{Cat} + FP_{Cat}} = \frac{64}{64 + 6}$$

		Predicted Label		
		Cat	Dog	Wolf
Actual Label	Cat	10	8	2
	Dog	5	30	15
	Wolf	1	4	15

Specificity for Multiple Classes

$$Specificity_{Cat} = \frac{TN_{Cat}}{TN_{Cat} + FP_{Cat}} = \frac{64}{64 + 6} = 0.9142$$

 $Specificity_{Cat}$ tells us that 91.42% of the pets that are **not cats** were not identified as **cats**.

		Predicted Label		
		Cat	Dog	Wolf
Actual Label	Cat	10	8	2
	Dog	5	30	15
	Wolf	1	4	15

Specificity for Multiple Classes

$$Specificity_{Cat} = \frac{TN_{Cat}}{TN_{Cat} + FP_{Cat}} = \frac{64}{64 + 6} = 0.9142$$



$Specificity_{Cat}$ tells us that 91.42% of the pets that are **not cats** were not identified as **cats**.

```
from sklearn.metrics import confusion_matrix  
  
cm = confusion_matrix(y_true, y_pred,  
                      labels=['Cat', 'Dog', 'Wolf'])
```

```
TP = np.diag(cm)  
FP = cm.sum(axis=0) - TP  
FN = cm.sum(axis=1) - TP  
TN = cm.sum() - (TP + FP + FN)
```

```
specificity = TN / (TN + FP)
```

```
[0.91428571 0.7 0.75714286]
```

Actual Label



Specificity for Multiple Classes

$$Specificity_{Cat} = \frac{TN_{Cat}}{TN_{Cat} + FP_{Cat}} = \frac{64}{64 + 6} = 0.9142$$

 $Specificity_{Cat}$ tells us that 91.42% of the pets that are **not cats** were not identified as **cats**.

```
from sklearn.metrics import confusion_matrix  
  
cm = confusion_matrix(y_true, y_pred,  
                      labels=['Cat', 'Dog', 'Wolf'])
```

```
TP = np.diag(cm)  
FP = cm.sum(axis=0) - TP  
FN = cm.sum(axis=1) - TP  
TN = cm.sum() - (TP + FP + FN)
```

```
specificity = TN / (TN + FP)
```

```
[0.91428571 0.7 0.75714286]
```

Actual Label



Single Specificity for Multiple Classes

macro

Calculate **metrics** for **each label** and **average** them.

For **c** classes, enumerated from **0** to **c-1**

$$Specificity_{macro} = \frac{\sum_{j=0}^{c-1} Specificity_j}{c}$$



Each class **contributes/weights equally** to the final metric since it **does not** take **class proportion (imbalance)** into account.
Thus, it is appropriate for **class-imbalanced datasets**.

```
from sklearn.metrics import confusion_matrix  
  
cm = confusion_matrix(y_true, y_pred,  
                      labels=['Cat', 'Dog', 'Wolf'])  
  
TP = np.diag(cm)  
FP = cm.sum(axis=0) - TP  
FN = cm.sum(axis=1) - TP  
TN = cm.sum() - (TP + FP + FN)  
  
specify = TN / (TN + FP)  
specify_macro = specify.mean()  
  
[0.91428571 0.7 0.75714286]  
0.7904761904761904
```

Single Specificity for Multiple Classes

weighted

Calculate metrics for each class, and find their **average weighted** by the **number of true samples of each class**.

Let c be the number of classes, enumerated from **0** to **c-1**,
 m the number samples, and m_j is the number of true samples of class j .

$$\text{Specificity}_{\text{weighted}} = \frac{\sum_{j=0}^{c-1} m_j * \text{Specificity}_j}{m}$$

weight (w_j) for class j

Single Specificity for Multiple Classes

weighted

Calculate metrics for each class, and find their **average weighted** by the **number of true samples of each class**.

Let c be the number of classes, enumerated from **0** to **c-1**,
 m the number samples, and m_j is the number of true samples of class j .

$$\text{Specificity}_{\text{weighted}} = \frac{\sum_{j=0}^{c-1} m_j * \text{Specificity}_j}{m}$$

weight (w_j) for class j



It **take into account class proportion (imbalance)** for computation.
Thus, it is **NOT appropriated** for **class-imbalanced datasets**.

Single Specificity for Multiple Classes

weighted

Calculate metrics for each class, and find their **average weighted** by the **number of true samples of each class**.

Let c be the number of classes, enumerated from **0** to **$c-1$** ,
 m the number samples, and m_j is the number of true samples of class j .

$$\text{Specificity}_{\text{weighted}} = \frac{\sum_{j=0}^{c-1} m_j * \text{Specificity}_j}{m}$$

weight (w_j) for class j



It take into account class proportion (imbalance) for computation.
Thus, it is **NOT appropriated** for class-imbalanced datasets.

```
...  
specify = TN / (TN + FP)  
  
# weighting  
class_sizes = cm.sum(axis=1)  
number_of_samples = cm.sum()  
weights = class_sizes / number_of_samples  
  
specify_weighted = np.sum(specify * weights)
```

```
[0.91428571 0.7 0.75714286]
```

```
[20 50 20]
```

```
90
```

```
[0.22222222 0.55555556 0.22222222]
```

```
0.7603174603174603
```

Exercise: Compute...

$$Sensitivity_{Dog} \quad Specificity_{Dog}$$
$$Sensitivity_{Wolf} \quad Specificity_{Wolf}$$

		Predicted Label		
		Cat	Dog	Wolf
Actual Label	Cat	10	8	2
	Dog	5	30	15
	Wolf	1	4	15

More Metrics

Precision

What is the **proportion** of **true positives** among all instances **classified** (**correctly** and **incorrectly**) as **positives**?
How precise is the **positive** classification?

“From all patients classified as **cancer**, how many (**proportion**) actually had **cancer**?”

$$Precision = \frac{TP}{TP + FP}$$

		Predicted Label	
		Negative	Positive
Actual Label	Negative	Healthy	Cancer
	Positive	TN	FP
Actual Label	Negative	Healthy	Cancer
	Positive	FN	TP

Precision

What is the **proportion** of **true positives** among all instances **classified** (**correctly** and **incorrectly**) as **positives**?
How precise is the **positive** classification?

“From all patients classified as **cancer**, how many (**proportion**) actually had **cancer**?”

$$Precision = \frac{TP}{TP + FP}$$

⚠️ **Negatives** are left out of the metric.

We can **cheat** by:

- Guess 1 **correct** guess as **positive** ($TP=1$)
- Guess all other samples (**positives** and **negatives**) as **negatives** ($FP=0$)
- **Precision = 1**

		Predicted Label	
		Negative	Positive
Actual	Positive	Healthy	Cancer
	Cancer	FN	TP

Precision

What is the **proportion** of **true positives** among all instances **classified** (**correctly** and **incorrectly**) as **positives**?
How precise is the **positive** classification?

“From all patients classified as **cancer**, how many (**proportion**) actually had **cancer**? ”

$$Precision = \frac{TP}{TP + FP}$$

`sklearn.metrics.precision_score`

```
from sklearn.metrics import precision_score  
precision_score(y_true, y_pred)
```

```
from sklearn.metrics import precision_score  
precision_score(y_true, y_pred, pos_label='Cancer')
```



By default, the **positive class** must have **value/label +1**.



We can **indicate** which is the **value/label** of the **positive class**.

$$Precision = \frac{TP}{TP + FP} =$$

		Predicted Label	
		Negative	Positive
Actual Label	Healthy	70	130
	Cancer	20	80

$$Precision = \frac{TP}{TP + FP} = \frac{80}{80 + 130} = 0.38$$

Only 38% of all people classified as having cancer actually had cancer.

		Predicted Label	
		Negative	Positive
Actual Label	Healthy	70	130
	Cancer	20	80

Precision for Multiple Classes

$$Precision_{Cat} = \frac{TP_{Cat}}{TP_{Cat} + FP_{Cat}} =$$



Similar to the other metrics.

		Predicted Label		
		Cat	Dog	Wolf
Actual Label	Cat	10	8	2
	Dog	5	30	15
	Wolf	1	4	15

Precision for Multiple Classes

$$Precision_{Cat} = \frac{TP_{Cat}}{TP_{Cat} + FP_{Cat}} = \frac{10}{10 + 6} = 0.625$$

 50% of pets classified as **Cat** are in fact **cats**.



Similar to the other metrics.

		Predicted Label		
		Cat	Dog	Wolf
Actual Label	Cat	10	8	2
	Dog	5	30	15
	Wolf	1	4	15

Precision for Multiple Classes

$$Precision_{Cat} = \frac{TP_{Cat}}{TP_{Cat} + FP_{Cat}} = \frac{10}{10 + 6} = 0.625$$



50% of pets classified as **Cat** are in fact **cats**.

```
from sklearn.metrics import precision_score  
precision_score(y_true, y_pred, average=None, labels=['Cat', 'Dog', 'Wolf'])  
  
array([0.625, 0.71428571, 0.46875])
```

	Cat	Dog	Wolf
Cat	10	8	2
Dog	5	30	15
Wolf	1	4	15

Precision for Multiple Classes

$$Precision_{Cat} = \frac{TP_{Cat}}{TP_{Cat} + FP_{Cat}} = \frac{10}{10 + 6} = 0.625$$



50% of pets classified as **Cat** are in fact **cats**.

```
from sklearn.metrics import precision_score
```

```
precision_score(y_true, y_pred)
```

We can return a **single “global” precision** by following the same strategies we have seen before!

		Cat	Dog	Wolf
Actual Label	Cat	10	8	2
	Dog	5	30	15
	Wolf	1	4	15

Single Precision for Multiple Classes

macro

Calculate **metrics** for **each label** and **average** them.

For **c** classes, enumerated from **0** to **c-1**

$$Precision_{macro} = \frac{\sum_{j=0}^{c-1} Precision_j}{c}$$



Each class **contributes/weights equally** to the final metric since it **does not** take **class proportion (imbalance)** into account.
Thus, it is appropriate for **class-imbalanced datasets**.

```
from sklearn.metrics import precision_score  
precision_score(y_true, y_pred, average='macro',  
                 labels=['Cat', 'Dog', 'Wolf'])
```

0.6026785714285715

Single Precision for Multiple Classes

macro

Calculate **metrics** for **each label** and **average** them.

For **c** classes, enumerated from **0** to **c-1**

$$Precision_{macro} = \frac{\sum_{j=0}^{c-1} Precision_j}{c}$$



Each class **contributes/weights equally** to the final metric since it **does not** take **class proportion (imbalance)** into account.
Thus, it is appropriate for **class-imbalanced datasets**.

micro

Calculate **metrics globally** by counting the total true positives, false negatives and false positives.

For **c** classes, enumerated from **0** to **c-1**, and **m** samples

$$Precision_{micro} = \frac{\sum_{j=0}^{c-1} TP_j}{m}$$

```
from sklearn.metrics import precision_score  
precision_score(y_true, y_pred, average='macro',  
                 labels=['Cat', 'Dog', 'Wolf'])
```

0.6026785714285715

```
from sklearn.metrics import precision_score  
precision_score(y_true, y_pred, average='micro',  
                 labels=['Cat', 'Dog', 'Wolf'])
```

0.6111111111111112

Single Precision for Multiple Classes

weighted

Calculate metrics for each class, and find their **average weighted** by the **number of true samples of each class**.

Let c be the number of classes, enumerated from **0** to **$c-1$** ,
 m the number samples, and m_j is the number of true samples of class j .

$$\text{Precision}_{\text{weighted}} = \frac{\sum_{j=0}^{c-1} m_j * \text{Precision}_j}{m}$$

weight (w_j) for class j



It take into account class proportion (imbalance) for computation.
Thus, it is **NOT appropriated** for class-imbalanced datasets.

```
from sklearn.metrics import precision_score  
precision_score(y_true, y_pred, average='weighted',  
                 labels=['Cat', 'Dog', 'Wolf'])
```

0.6398809523809524

Exercise: Compute...

$Precision_{Dog}$

$Precision_{Wolf}$

		Predicted Label		
		Cat	Dog	Wolf
Actual Label	Cat	10	8	2
	Dog	5	30	15
	Wolf	1	4	15

Accuracy

What was the (overall) classification hit rate?

$$\text{Accuracy} = \frac{TP + TN}{TN + FN + FP + TP}$$

		Predicted Label	
		Negative	Positive
Actual Label	Negative	Healthy	Cancer
	Positive	TN	FP
	Negative	FN	TP

Accuracy

What was the (overall) classification hit rate?

$$\text{Accuracy} = \frac{TP + TN}{TN + FN + FP + TP}$$

Issue with **biased dataset (imbalance)**:

Easy to **cheat** the score:

- Consider that only **1%** are **positives**
- Guess everything as **negative**
- Achieve **99% accuracy**

It **take into account class proportion (imbalance)** for computation.

Thus, it is **NOT appropriated** for **class-imbalanced datasets**.

		Predicted Label	
		Negative	Positive
Actual	Positive	Healthy	Cancer
	Negative	TN	FP
	Cancer	FN	TP

Accuracy

What was the (overall) **classification hit rate**?

$$\text{Accuracy} = \frac{TP + TN}{TN + FN + FP + TP}$$



The same idea is applied to **extend it to multiple classes**.

Issue with **biased dataset (imbalance)**:

Easy to **cheat** the score:

- Consider that only **1%** are **positives**
- Guess everything as **negative**
- Achieve **99% accuracy**

⚠ It **take into account class proportion (imbalance)** for computation.

Thus, it is **NOT appropriated** for **class-imbalanced datasets**.

		Predicted Label	
		Negative	Positive
Actual	Positive	Healthy	Cancer
		TN	FP
	Cancer	FN	TP

Accuracy

What was the (overall) classification hit rate?

$$Accuracy = \frac{TP + TN}{TN + FN + FP + TP}$$

`sklearn.metrics.accuracy_score`

```
from sklearn.metrics import accuracy_score  
accuracy_score(y_true, y_pred)
```



We **do not** need to specify which is the **positive** (or **negative**) class, since **accuracy** takes into account **all metrics** from the **confusion matrix**.

$$Accuracy = \frac{TP + TN}{TN + FN + FP + TP}$$

		Predicted Label	
		Negative	Positive
Actual Label	Healthy	Healthy	Cancer
	Healthy	70	130
Positive		Cancer	20
			80

$$Accuracy = \frac{70 + 80}{70 + 20 + 130 + 80} = \frac{150}{300} = 0.5$$

The overall classification hit rate is 50%.

		Predicted Label	
		Negative	Positive
Actual Label	Healthy	Healthy	Cancer
	Healthy	70	130
		Cancer	20
			80

Accuracy for Multiple Classes

For **c** classes, enumerated from **0** to **c-1**, and **m** samples

$$Accuracy = \frac{\sum_{j=0}^{c-1} TP_j}{m} = \frac{TP_{Cat} + TP_{Dog} + TP_{Wolf}}{m}$$

Sum of all values from
the confusion matrix

		Predicted Label		
		Cat	Dog	Wolf
Actual Label	Cat	10	8	2
	Dog	5	30	15
	Wolf	1	4	15

Accuracy for Multiple Classes

For **c** classes, enumerated from **0** to **c-1**, and **m** samples

$$Accuracy = \frac{\sum_{j=0}^{c-1} TP_j}{m} = \frac{TP_{Cat} + TP_{Dog} + TP_{Wolf}}{m}$$

Sum of all values from
the confusion matrix

$$Accuracy = \frac{10 + 30 + 15}{90} = 0.6111$$

		Predicted Label		
		Cat	Dog	Wolf
Actual Label	Cat	10	8	2
	Dog	5	30	15
	Wolf	1	4	15

Accuracy for Multiple Classes

For **c** classes, enumerated from **0** to **c-1**, and **m** samples

$$Accuracy = \frac{\sum_{j=0}^{c-1} TP_j}{m} = \frac{TP_{Cat} + TP_{Dog} + TP_{Wolf}}{m}$$

Sum of all values from the confusion matrix

$$Accuracy = \frac{10 + 30 + 15}{90} = 0.6111$$

		Predicted Label		
		Cat	Dog	Wolf
Actual Label	Cat	10	8	2
	Dog	30	15	
	Wolf	1	4	15



It **take into account class proportion (imbalance)** for computation.
Thus, it is **NOT appropriated for class-imbalanced datasets.**

Balanced Accuracy

Adapted version of **accuracy** that deals with **class-imbalanced datasets**.

Each class **contributes/weights equally** to the final metric.

It is defined as the **average of sensitivity/recall** obtained on each class.



It is equal to ***Recall_{macro}***.

Balanced Accuracy

Adapted version of **accuracy** that deals with **class-imbalanced datasets**.
Each class **contributes/weights equally** to the final metric.

It is defined as the **average of sensitivity/recall** obtained on each class.



It is equal to ***Recall_{macro}***.

Binary case

$$\text{BalancedAccuracy} = \frac{\text{Sensitivity} + \text{Specificity}}{(2)}$$

“sensitivity/recall” of
the **negative class**

Two classes.

Balanced Accuracy

Adapted version of **accuracy** that deals with **class-imbalanced datasets**.
Each class **contributes/weights equally** to the final metric.

It is defined as the **average of sensitivity/recall** obtained on each class.



It is equal to ***Recall_{macro}***.

Binary case

$$\text{BalancedAccuracy} = \frac{\text{Sensitivity} + \text{Specificity}}{2}$$

“sensitivity/recall” of
the **negative class**

Two classes.

Multiclass

For **c** classes, enumerated from **0** to **c-1**

$$\text{BalancedAccuracy} = \frac{\sum_{j=0}^{c-1} \text{Recall}_j}{c}$$

Balanced Accuracy

Adapted version of **accuracy** that deals with **class-imbalanced datasets**.
Each class **contributes/weights equally** to the final metric.

It is defined as the **average of sensitivity/recall** obtained on each class.



It is equal to ***Recall_{macro}***.

Binary case

$$\text{BalancedAccuracy} = \frac{\text{Sensitivity} + \text{Specificity}}{2}$$

“sensitivity/recall” of
the **negative class**

Two classes.

Multiclass

For **c** classes, enumerated from **0** to **c-1**

$$\text{BalancedAccuracy} = \frac{\sum_{j=0}^{c-1} \text{Recall}_j}{c}$$

```
from sklearn.metrics import balanced_accuracy_score  
balanced_accuracy_score(y_true, y_pred, adjusted=False)
```

$$BalancedAccuracy = \frac{Sensitivity + Specificity}{2}$$

		Predicted Label	
		Negative	Positive
Actual Label	Healthy	Healthy	Cancer
	Healthy	70	130
Positive		Cancer	20
Negative			80

$$BalancedAccuracy = \frac{0.8 + 0.35}{2} = 0.575$$

 The **class-balanced** overall classification hit rate is **57.5%**.

		Predicted Label	
		Negative	Positive
Actual Label Positive	Healthy	Healthy	Cancer
	Healthy	70	130
Negative		Cancer	20
			80

Balanced Accuracy for Multiple Classes

For c classes, enumerated from 0 to $c-1$

$$\text{BalancedAccuracy} = \frac{\sum_{j=0}^{c-1} \text{Recall}_j}{c} = \frac{\text{Recall}_{\text{Cat}} + \text{Recall}_{\text{Dog}} + \text{Recall}_{\text{Wolf}}}{3}$$

		Predicted Label		
		Cat	Dog	Wolf
Actual Label	Cat	10	8	2
	Dog	5	30	15
	Wolf	1	4	15

Balanced Accuracy for Multiple Classes

For c classes, enumerated from 0 to $c-1$

$$\text{BalancedAccuracy} = \frac{\sum_{j=0}^{c-1} \text{Recall}_j}{c} = \frac{0.5 + 0.6 + 0.75}{3} = 0.616666$$

 The **class-balanced** overall classification hit rate is **61.666%**.

		Predicted Label		
		Cat	Dog	Wolf
Actual Label	Cat	10	8	2
	Dog	5	30	15
	Wolf	1	4	15

Class Balanced Accuracy (CBA)

A version proposed by [Mosley et al.](#), “A balanced approach to the multi-class imbalance problem”, IJCV 2010.

- Section 3.2, Equations 3.1 – 3.4

For **c** classes, enumerated from **0** to **c-1**

$$CBA = \frac{\sum_{j=0}^{c-1} \min(Recall_j, Precision_j)}{c}$$

Class Balanced Accuracy (CBA)

A version proposed by [Mosley et al.](#), “A balanced approach to the multi-class imbalance problem”, IJCV 2010.

- Section 3.2, Equations 3.1 – 3.4

For c classes, enumerated from 0 to $c-1$

$$CBA = \frac{\sum_{j=0}^{c-1} \min(Recall_j, Precision_j)}{c}$$



It is a **more rigid metric** than **Balanced Accuracy**, penalizing **class imbalance** and classification errors **more strongly**.

Class Balanced Accuracy (CBA)

A version proposed by [Mosley et al.](#), “A balanced approach to the multi-class imbalance problem”, IJCV 2010.

- Section 3.2, Equations 3.1 – 3.4

For c classes, enumerated from 0 to $c-1$

$$CBA = \frac{\sum_{j=0}^{c-1} \min(Recall_j, Precision_j)}{c}$$



It is a **more rigid metric** than **Balanced Accuracy**, penalizing **class imbalance** and classification errors **more strongly**.

```
def class_balanced_accuracy(y_true, y_pred):
    labels = list(set(y_true))

    recall = recall_score(y_true, y_pred, average=None, labels=labels)
    precision = precision_score(y_true, y_pred, average=None, labels=labels)

    cba = np.minimum(recall, precision).mean()
    return cba

class_balanced_accuracy(y_true, y_pred)
```

For c classes, enumerated from 0 to $c-1$

$$CBA = \frac{\sum_{j=0}^{c-1} \min(Recall_j, Precision_j)}{c}$$

		Predicted Label	
		Negative	Positive
Actual Label	Healthy	Healthy	Cancer
	Healthy	70	130
		Cancer	80
		20	

$$CBA = \frac{\min(Recall_{Healthy}, Precision_{Healthy}) + \min(Recall_{Cancer}, Precision_{Cancer})}{2}$$

		Predicted Label	
		Negative	Positive
		Healthy	Cancer
Actual Label Positive	Healthy	70	130
	Cancer	20	80

$$CBA = \frac{\min(0.35, 0.7778) + \min(0.8, 0.381)}{2}$$

		Predicted Label	
		Negative	Positive
Actual Label	Healthy	Healthy	Cancer
	Healthy	70	130
		Cancer	80
		20	

$$CBA = \frac{\min(0.35, 0.7778) + \min(0.8, 0.381)}{2} = \frac{0.35 + 0.381}{2} = 0.3655$$

		Predicted Label	
		Negative	Positive
		Healthy	Cancer
Actual Label Positive	Healthy	70	130
	Cancer	20	80

Exercise: Compute...

For c classes, enumerated from 0 to $c-1$

$$CBA = \frac{\sum_{j=0}^{c-1} \min(Recall_j, Precision_j)}{c}$$

$$CBA = \frac{\min(Precision_{Cat}, Precision_{Cat}) + \min(Precision_{Dog}, Precision_{Dog}) + \min(Precision_{Wolf}, Precision_{Wolf})}{3}$$

		Predicted Label		
		Cat	Dog	Wolf
Actual Label	Cat	10	8	2
	Dog	5	30	15
	Wolf	1	4	15

F1-Score

How **good** and **complete** are the predictions?
Middle ground between **precision** and **recall**.

$$F1 = 2 * \frac{Precision * Recall}{Precision + Recall}$$



It solves the **accuracy cheating** caused by **class imbalance**.



The same idea is applied to **extend it to multiple classes**.

		Predicted Label		
		Negative	Positive	
Actual Label	Negative	Healthy	Cancer	
	Positive	Healthy	TN	FP
	Positive	Cancer	FN	TP

F1-Score

How **good** and **complete** are the predictions?
Middle ground between **precision** and **recall**.

$$F1 = 2 * \frac{Precision * Recall}{Precision + Recall}$$

`sklearn.metrics.f1_score`

```
from sklearn.metrics import f1_score  
f1_score(y_true, y_pred)
```



By default, the **positive class** must have **value/label +1**.

```
from sklearn.metrics import f1_score  
f1_score(y_true, y_pred, pos_label='Cancer')
```



We can **indicate** which is the **value/label** of the **positive class**.

$$F1 = 2 * \frac{Precision * Recall}{Precision + Recall}$$

		Predicted Label	
		Negative	Positive
Actual Label	Healthy	Healthy	Cancer
	Healthy	70	130
Positive		Cancer	20
			80

$$F1 = 2 * \frac{Precision * Recall}{Precision + Recall} = 2 * \frac{0.381 * 0.8}{0.381 + 0.8} = 0.5161$$

		Predicted Label	
		Negative	Positive
Actual Label	Healthy	Healthy	Cancer
	Healthy	70	130
		Cancer	80
		20	

F1-Score for Multiple Classes

$$F1_{Cat} = 2 * \frac{Precision_{Cat} * Recall_{Cat}}{Precision_{Cat} + Recall_{Cat}}$$

		Predicted Label		
		Cat	Dog	Wolf
Actual Label	Cat	10	8	2
	Dog	5	30	15
	Wolf	1	4	15

F1-Score for Multiple Classes

$$F1_{Cat} = 2 * \frac{Precision_{Cat} * Recall_{Cat}}{Precision_{Cat} + Recall_{Cat}} = 2 * \frac{0.625 * 0.5}{0.625 + 0.5} = 0.555$$

		Predicted Label		
		Cat	Dog	Wolf
Actual Label	Cat	10	8	2
	Dog	5	30	15
	Wolf	1	4	15

F1-Score for Multiple Classes

$$F1_{Cat} = 2 * \frac{Precision_{Cat} * Recall_{Cat}}{Precision_{Cat} + Recall_{Cat}} = 2 * \frac{0.625 * 0.5}{0.625 + 0.5} = 0.555$$

```
from sklearn.metrics import f1_score
f1_score(y_true, y_pred, average=None, labels=['Cat', 'Dog', 'Wolf'])

array([0.55555556, 0.65217391, 0.57692308])
```

	Cat	Dog	Wolf
Cat	10	8	2
Dog	5	30	15
Wolf	1	4	15

F1-Score for Multiple Classes

$$F1_{Cat} = 2 * \frac{Precision_{Cat} * Recall_{Cat}}{Precision_{Cat} + Recall_{Cat}} = 2 * \frac{0.625 * 0.5}{0.625 + 0.5} = 0.555$$

```
from sklearn.metrics import f1_score
f1_score(array([0, 1, 2, 0, 1, 2, 0, 1, 2]), array([0, 0, 1, 1, 1, 2, 2, 2, 2]), average='macro')
```

We can return a **single “global” F1-Score** by following the same strategies we have seen before!

		Cat	Dog	Wolf
Actual Label	Cat	10	8	2
	Dog	5	30	15
	Wolf	1	4	15

Single F1-Score for Multiple Classes

macro

Calculate **metrics** for **each label** and **average** them.

For **c** classes, enumerated from **0** to **c-1**

$$F1_{macro} = \frac{\sum_{j=0}^{c-1} F1_j}{c}$$



Each class **contributes/weights equally** to the final metric since it **does not** take **class proportion (imbalance)** into account.
Thus, it is appropriate for **class-imbalanced datasets**.

```
from sklearn.metrics import f1_score  
f1_score(y_true, y_pred, average='macro',  
          labels=['Cat', 'Dog', 'Wolf'])
```

0.5948841818407036

Single F1-Score for Multiple Classes

macro

Calculate **metrics** for **each label** and **average** them.

For **c** classes, enumerated from **0** to **c-1**

$$F1_{macro} = \frac{\sum_{j=0}^{c-1} F1_j}{c}$$



Each class **contributes/weights equally** to the final metric since it **does not** take **class proportion (imbalance)** into account.
Thus, it is appropriate for **class-imbalanced datasets**.

micro

Calculate **metrics globally** by counting the total true positives, false negatives and false positives.

For **c** classes, enumerated from **0** to **c-1**, and **m** samples

$$F1_{micro} = \frac{\sum_{j=0}^{c-1} TP_j}{m}$$

```
from sklearn.metrics import f1_score  
f1_score(y_true, y_pred, average='macro',  
          labels=['Cat', 'Dog', 'Wolf'])
```

0.5948841818407036

```
from sklearn.metrics import f1_score  
f1_score(y_true, y_pred, average='micro',  
          labels=['Cat', 'Dog', 'Wolf'])
```

0.6111111111111112

Single F1-Score for Multiple Classes

weighted

Calculate metrics for each class, and find their **average weighted** by the **number of true samples of each class**.

Let c be the number of classes, enumerated from **0** to **c-1**,
 m the number samples, and m_j is the number of true samples of class j .

$$F1_{weighted} = \frac{\sum_{j=0}^{c-1} m_j * F1_j}{m}$$

weight (w_j) for class j



It take into account class proportion (imbalance) for computation.
Thus, it is NOT appropriated for class-imbalanced datasets.

```
from sklearn.metrics import f1_score  
f1_score(y_true, y_pred, average='weighted',  
          labels=['Cat', 'Dog', 'Wolf'])
```

0.6139807589082951

Putting it all together

Binary problem

```
from sklearn.metrics import classification_report  
print(classification_report(y_true, y_pred, labels=['Healthy', 'Cancer'], digits=4))
```

	precision	recall	f1-score	support
Healthy	0.7778	0.3500	0.4828	200
Cancer	0.3810	0.8000	0.5161	100
accuracy			0.5000	300
macro avg	0.5794	0.5750	0.4994	300
weighted avg	0.6455	0.5000	0.4939	300

Putting it all together

Binary problem

Provides a **report** with some
classification metrics.

```
from sklearn.metrics import classification_report
print(classification_report(y_true, y_pred, labels=['Healthy', 'Cancer'], digits=4))
```

	precision	recall	f1-score	support
Healthy	0.7778	0.3500	0.4828	200
Cancer	0.3810	0.8000	0.5161	100
accuracy			0.5000	300
macro avg	0.5794	0.5750	0.4994	300
weighted avg	0.6455	0.5000	0.4939	300

Putting it all together

Binary problem

Optional: it is just to keep a **given desired order** in the report.

```
from sklearn.metrics import classification_report  
print(classification_report(y_true, y_pred, labels=['Healthy', 'Cancer'], digits=4))
```

	precision	recall	f1-score	support
Healthy	0.7778	0.3500	0.4828	200
Cancer	0.3810	0.8000	0.5161	100
accuracy			0.5000	300
macro avg	0.5794	0.5750	0.4994	300
weighted avg	0.6455	0.5000	0.4939	300

Putting it all together

Binary problem

```
from sklearn.metrics import classification_report  
print(classification_report(y_true, y_pred, labels=['Healthy', 'Cancer'], digits=4))
```

	precision	recall	f1-score	support	Number of samples of each class.
Healthy	0.7778	0.3500	0.4828	200	
Cancer	0.3810	0.8000	0.5161	100	
accuracy			0.5000	300	
macro avg	0.5794	0.5750	0.4994	300	
weighted avg	0.6455	0.5000	0.4939	300	

Putting it all together

Binary problem

```
from sklearn.metrics import classification_report  
print(classification_report(y_true, y_pred, labels=['Healthy', 'Cancer'], digits=4))
```

	precision	recall	f1-score	support
Healthy	0.7778	0.3500	0.4828	200
Cancer	0.3810	0.8000	0.5161	100
accuracy			0.5000	300
macro avg	0.5794	0.5750	0.4994	300
weighted avg	0.6455	0.5000	0.4939	300

Balanced Accuracy

Putting it all together

Multiclass problem

```
from sklearn.metrics import classification_report  
print(classification_report(y_true, y_pred, labels=['Cat', 'Dog', 'Wolf'], digits=4))
```

	precision	recall	f1-score	support
Cat	0.6250	0.5000	0.5556	20
Dog	0.7143	0.6000	0.6522	50
Wolf	0.4688	0.7500	0.5769	20
accuracy			0.6111	90
macro avg	0.6027	0.6167	0.5949	90
weighted avg	0.6399	0.6111	0.6140	90

Accuracy vs F1-Score

Accuracy

F1-Score

Accuracy vs F1-Score

Accuracy

1. Used when the **True Positives** and **True Negatives** are **more important** (overall hit rate);

F1-Score

1. Used when the **False Negatives** and **False Positives** are **crucial** (classification errors);

Accuracy vs F1-Score

Accuracy

1. Used when the **True Positives** and **True Negatives** are **more important** (overall hit rate);
2. Suitable for **(roughly) class-balanced** datasets;

F1-Score

1. Used when the **False Negatives** and **False Positives** are **crucial** (classification errors);
2. Works fine for **class-imbalanced** datasets;

Accuracy vs F1-Score

Accuracy

1. Used when the **True Positives** and **True Negatives** are **more important** (overall hit rate);
2. Suitable for **(roughly) class-balanced** datasets;

F1-Score

1. Used when the **False Negatives** and **False Positives** are **crucial** (classification errors);
2. Works fine for **class-imbalanced** datasets;
3. One of the **most used** classification metric;

Accuracy vs F1-Score

Accuracy

1. Used when the **True Positives** and **True Negatives** are **more important** (overall hit rate);
2. Suitable for **(roughly) class-balanced** datasets;
3. Its **balanced version** (**Balanced Accuracy - BA**) deals with **class-imbalance** while keeping the goal (1);

F1-Score

1. Used when the **False Negatives** and **False Positives** are **crucial** (classification errors);
2. Works fine for **class-imbalanced** datasets;
3. One of the **most used** classification metric;

Accuracy vs F1-Score

Accuracy

1. Used when the **True Positives** and **True Negatives** are **more important** (overall hit rate);
2. Suitable for **(roughly) class-balanced** datasets;
3. Its **balanced version** (**Balanced Accuracy - BA**) deals with **class-imbalance** while keeping the goal (1);
4. The **balanced version** by Mosley et al. (**Class Balanced Accuracy - CBA**) is **more rigid** than **BA** when **strongly penalizing** class imbalance and classification errors;

F1-Score

1. Used when the **False Negatives** and **False Positives** are **crucial** (classification errors);
2. Works fine for **class-imbalanced** datasets;
3. One of the **most used** classification metric;

Accuracy vs F1-Score

Accuracy

1. Used when the **True Positives** and **True Negatives** are **more important** (overall hit rate);
2. Suitable for **(roughly) class-balanced** datasets;
3. Its **balanced version** (**Balanced Accuracy - BA**) deals with **class-imbalance** while keeping the goal (1);
4. The **balanced version** by Mosley et al. (**Class Balanced Accuracy - CBA**) is **more rigid** than **BA** when **strongly penalizing** class imbalance and classification errors;
5. **CBA** is more interesting for **rigorous experiments** where **True Positives** and **True negatives** are **more important** ;

F1-Score

1. Used when the **False Negatives** and **False Positives** are **crucial** (classification errors);
2. Works fine for **class-imbalanced** datasets;
3. One of the **most used** classification metric;

Putting it all together

Binary problem

	Sensitivity/ Recall	Specificity	Precision	Accuracy	Balanced Accuracy	Class Balanced Accuracy	F1-Score
Classifier 1 (KNN)	0.8	0.35	0.38	0.5	0.575	0.365	0.516
Classifier 2 (Log. Regression)	0.5	0.4	0.29	0.433	0.45	0.347	0.37

Putting it all together

Binary problem

	Sensitivity/ Recall	Specificity	Precision	Accuracy	Balanced Accuracy	Class Balanced Accuracy	F1-Score
Classifier 1 (KNN)	0.8	0.35	0.38	0.5	0.575	0.365	0.516
Classifier 2 (Log. Regression)	0.5	0.4	0.29	0.433	0.45	0.347	0.37

We should pick **KNN** since it is **consistently better** than Log. Regression for **Balanced** and **Class Balanced Accuracy**, and **F1-Score**

Putting it all together

Multiclass

	Sensitivity/ Recall	Specificity	Precision	Accuracy	Balanced Accuracy	Class Balanced Accuracy	F1-Score
Cat	0.5	0.9142	0.625	--	--		0.555
Dog	0.6	0.7	0.7143	--	--		0.652
Wolf	0.75	0.757	0.4688	--	--		0.576
Macro	0.6167	0.79	0.6	0.6111	0.6167	0.522	0.594

Putting it all together

Multiclass

	Sensitivity/ Recall	Specificity	Precision	Accuracy	Balanced Accuracy	Class Balanced Accuracy	F1-Score
Cat	0.5	0.9142	0.625	--	--		0.555
Dog	0.6	0.7	0.7143	--	--		0.652
Wolf	0.75	0.757	0.4688	--	--		0.576
Macro	0.6167	0.79	0.6	0.6111	0.6167	0.522	0.594

The classifier is **not effective** when presenting
poor evaluation scores.

Putting it all together

Multiclass

	Sensitivity/ Recall	Specificity	Precision	Accuracy	Balanced Accuracy	Class Balanced Accuracy	F1-Score
Cat	0.5	0.9142	0.625	--	--		0.555
Dog	0.6	0.7	0.7143	--	--		0.652
Wolf	0.75	0.757	0.4688	--	--		0.576
Macro	0.6167	0.79	0.6	0.6111	0.6167	0.522	0.594

It considerably fails when classifying **Cats** and **Dogs**:
low sensitivity results, for example, in many cats not classified as cats.

Putting it all together

Multiclass

	Sensitivity/ Recall	Specificity	Precision	Accuracy	Balanced Accuracy	Class Balanced Accuracy	F1-Score
Cat	0.5	0.9142	0.625	--	--		0.555
Dog	0.6	0.7	0.7143	--	--		0.652
Wolf	0.75	0.757	0.4688	--	--		0.576
Macro	0.6167	0.79	0.6	0.6111	0.6167	0.522	0.594

The **high specificity** for Cats indicates that the classifier is correctly classifying non-cats as non-cats.

However, this does not necessarily mean that the classifier is not making mistakes.

For example, a dog (non-cat) may have been classified as a wolf (non-cat) and vice versa.

Putting it all together

Multiclass

	Sensitivity/ Recall	Specificity	Precision	Accuracy	Balanced Accuracy	Class Balanced Accuracy	F1-Score
Cat	0.5	0.9142	0.625	--	--		0.555
Dog	0.6	0.7	0.7143	--	--		0.652
Wolf	0.75	0.757	0.4688	--	--		0.576
Macro	0.6167	0.79	0.6	0.6111	0.6167	0.522	0.594

Although most **Wolfs** in the dataset were **correctly** classified as **Wolfs**...

Putting it all together

Multiclass

	Sensitivity/ Recall	Specificity	Precision	Accuracy	Balanced Accuracy	Class Balanced Accuracy	F1-Score
Cat	0.5	0.9142	0.625	--	--		0.555
Dog	0.6	0.7	0.7143	--	--		0.652
Wolf	0.75	0.757	0.4688	--	--		0.576
Macro	0.6167	0.79	0.6	0.6111	0.6167	0.522	0.594

Although most **Wolfs** in the dataset were **correctly classified as Wolfs**, the classifier tends to classify **most pets as Wolfs** (see the **low precision**).

Other Metrics

- ROC curve and AUC
- Fbeta-Score
- G-Mean
- Jaccard score
- Matthews correlation coefficient
- Cohen's kappa
- Brier score

Aprendizado de Máquina e Reconhecimento de Padrões 2021.2



Evaluation Metrics for Classification

Prof. Samuel Martins (Samuka)
samuel.martins@ifsp.edu.br

